

Introduction to Dask

PARALLEL PROGRAMMING WITH DASK IN PYTHON



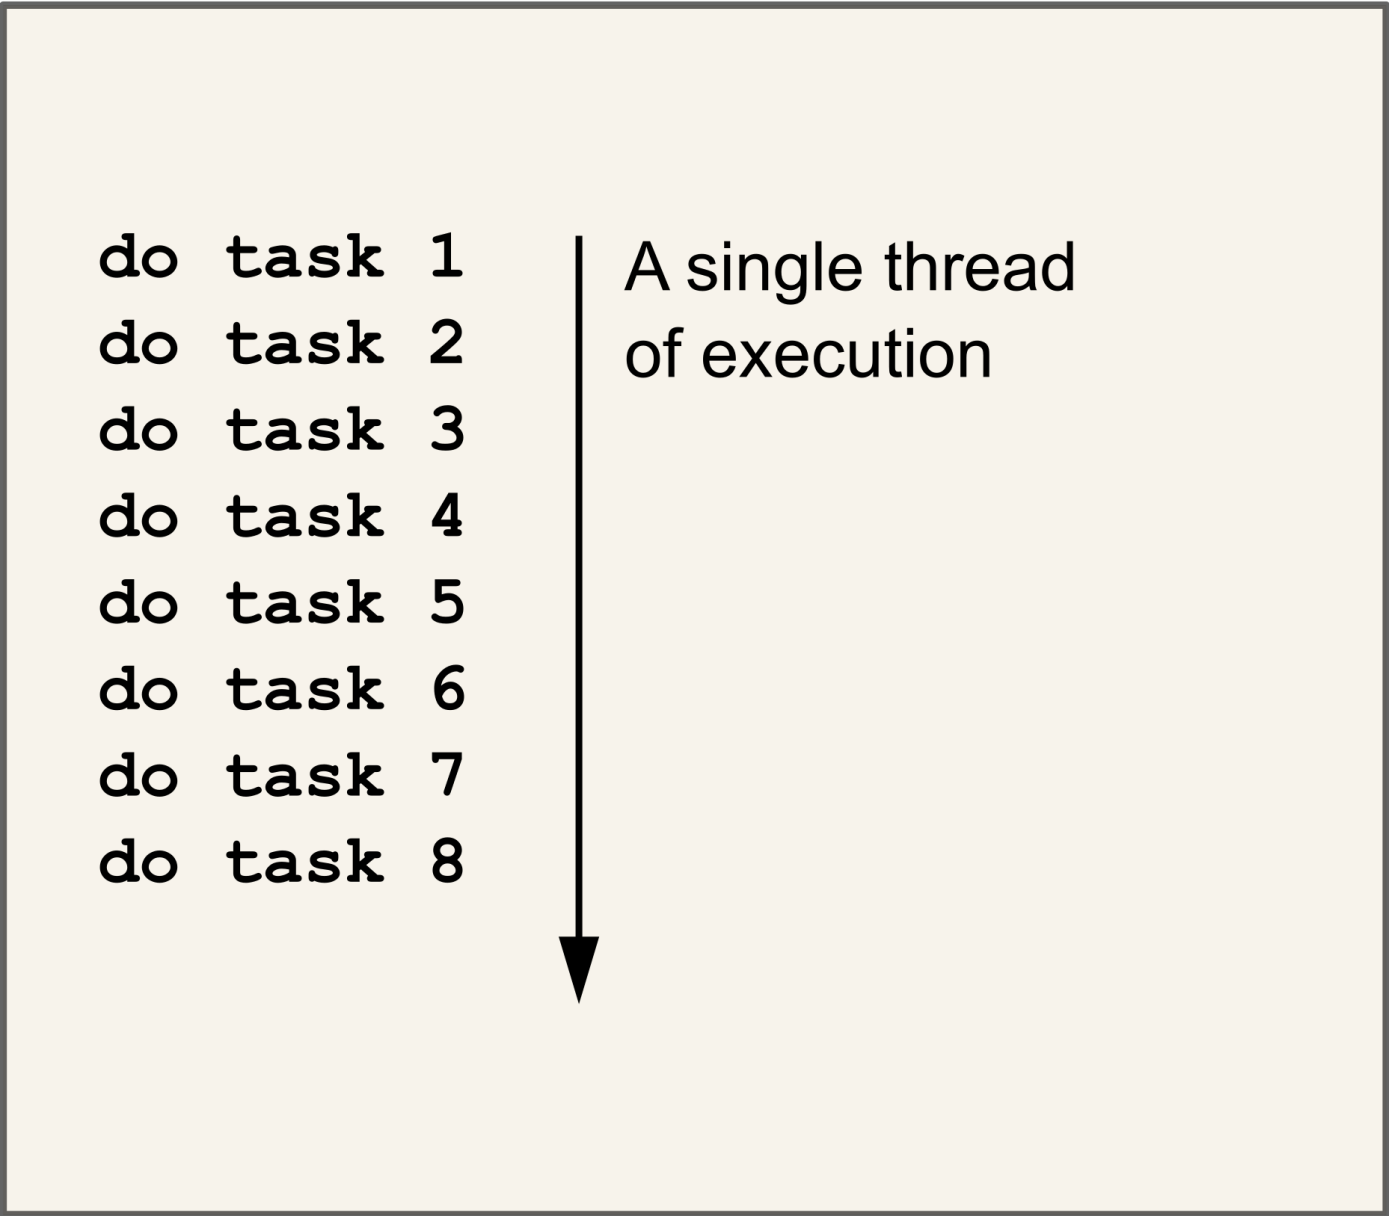
James Fulton

Climate Informatics Researcher

Speeding up computations using multiple cores

- Computers have multiple cores
- Code needs to be written to use them
- The Dask package can be used to do this
- Complete our computations faster

Concurrent programming



A diagram illustrating a single thread of execution. It features a light beige rectangular box with a thin black border. Inside the box, on the left side, is a list of eight tasks: 'do task 1', 'do task 2', 'do task 3', 'do task 4', 'do task 5', 'do task 6', 'do task 7', and 'do task 8'. To the right of this list, the text 'A single thread of execution' is written. A vertical black arrow points downwards from the text to the bottom of the box, indicating the sequential flow of the tasks.


```
do task 1
do task 2
do task 3
do task 4
do task 5
do task 6
do task 7
do task 8
```

A single thread
of execution


Multi-threading

Two threads of execution


do task 1
do task 2
do task 3
do task 4




do task 5
do task 6
do task 7
do task 8




Multi-threading



Multi-threading

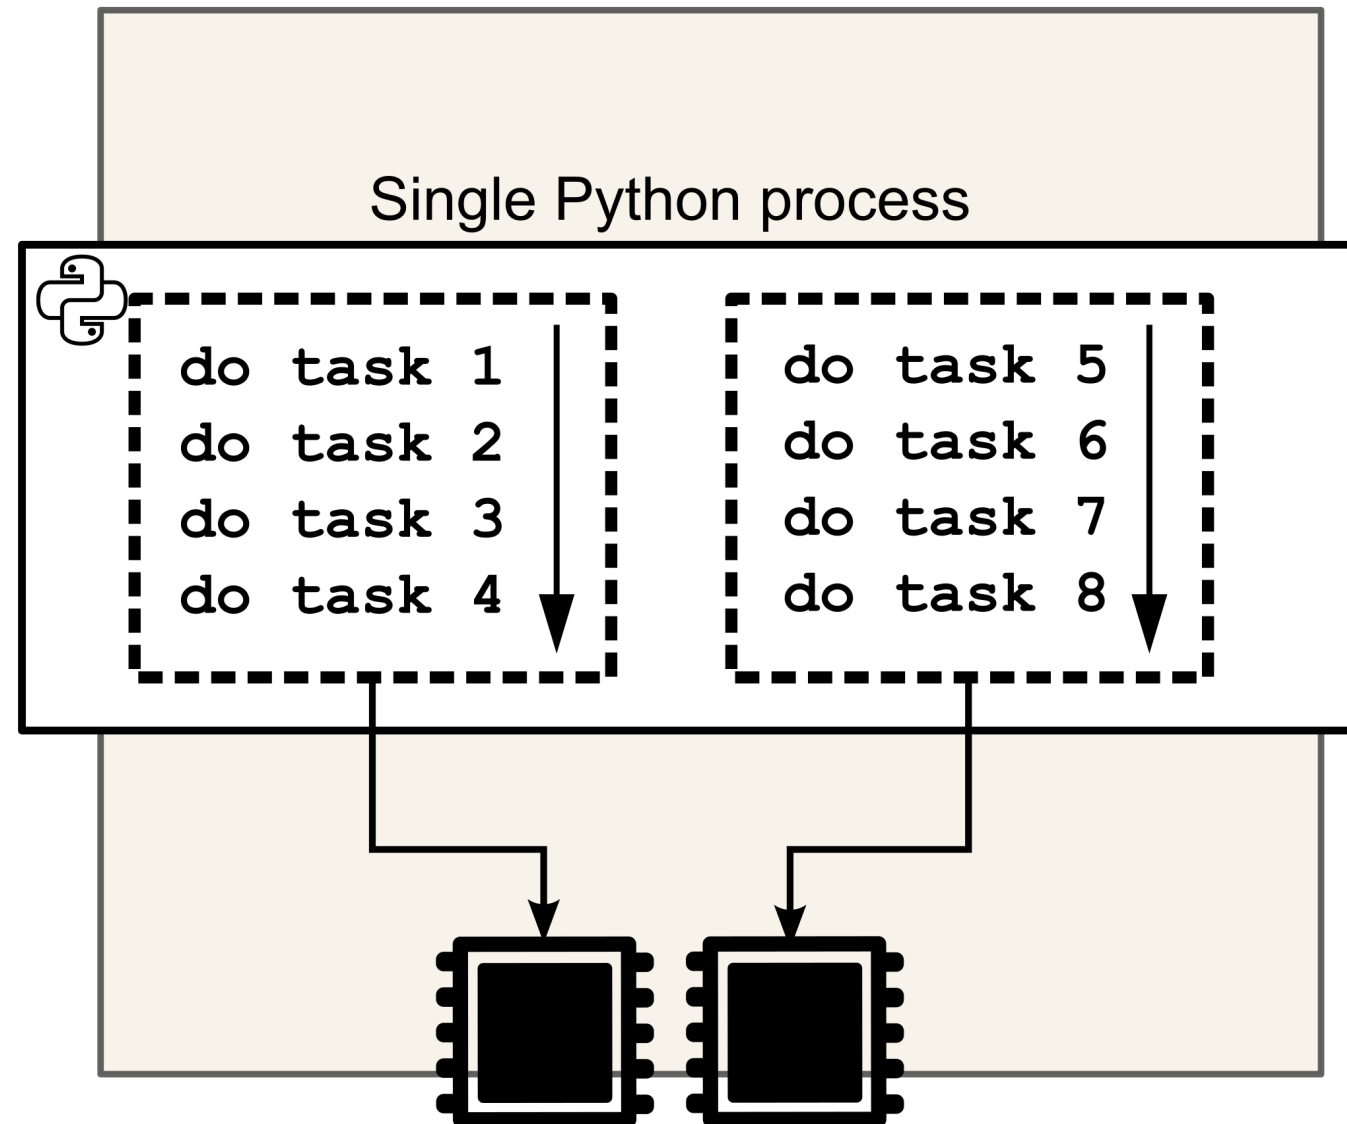


Parallel processing

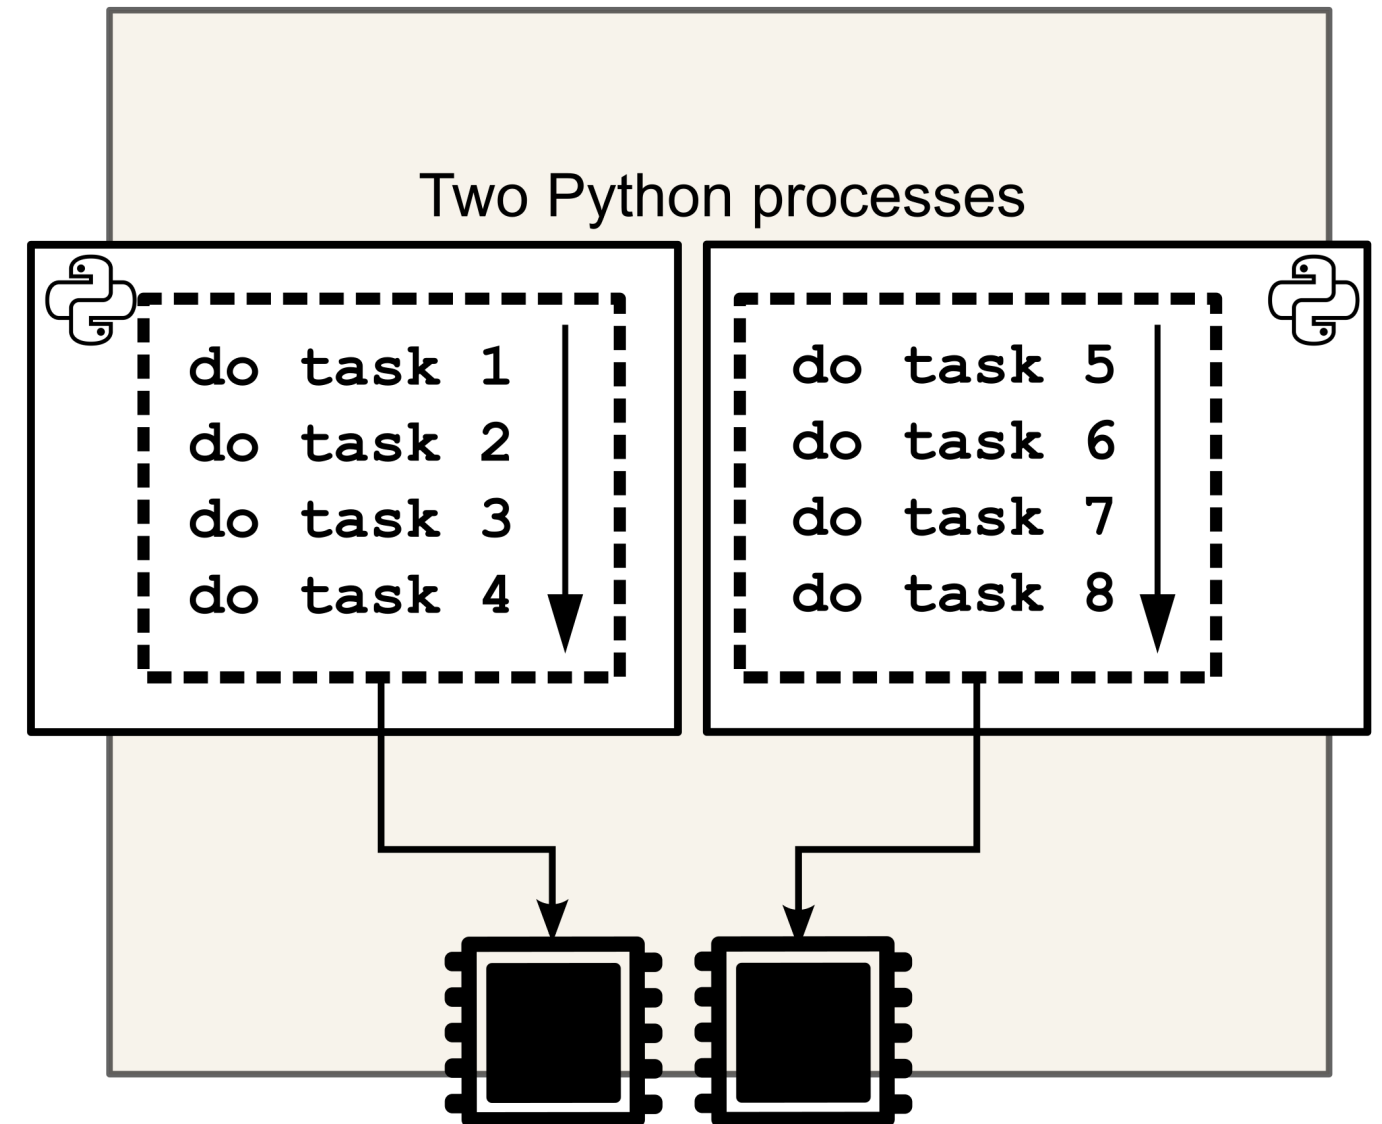


Parallel programming

Multi-threading



Parallel processing



Lazy evaluation

- Computations are not run until the moment the result is needed
- The steps required to compute the result are stored for later
- Dask splits the tasks between threads or processes

Dask delayed

```
from dask import delayed

def my_square_function(x):
    return x**2

# Create delayed version of above function
delayed_square_function = delayed(my_square_function)
```

Dask delayed

```
from dask import delayed

def my_square_function(x):
    return x**2

# Create delayed version of above function
delayed_square_function = delayed(my_square_function)

# Use the delayed function with input 4
delayed_result = delayed_square_function(4)

# Print the delayed answer
print(delayed_result)
```

```
Delayed('my_square_function-7f71b132-70a9-457a-aa52-604e8c34f8a7')
```

Dask delayed

```
from dask import delayed

def my_square_function(x):
    return x**2

# Delay and use function
delayed_result = delayed(my_square_function)(4)

print(delayed_result)
```

```
Delayed('my_square_function-7f71b132-70a9-457a-aa52-604e8c34f8a7')
```

Computing the answer

```
from dask import delayed

def my_square_function(x):
    return x**2

delayed_result = delayed(my_square_function)(4)

real_result = delayed_result.compute() # <- This line is where the calculation happens

# Print the answer
print(real_result)
```

16

Using operations on delayed objects

```
delayed_result1 = delayed(my_square_function)(4)

# Math operations return delayed object
delayed_result2 = (4 + delayed_result1) * 5

print(delayed_result2.compute())
```

100

Lazy evaluation

```
x_list = [30, 85, 14, 12, 27, 62, 89, 15, 78, 0]

sum_of_squares = 0

for x in x_list:
    # Square and add numbers
    sum_of_squares += delayed(my_square_function)(x)
```

Lazy evaluation

```
x_list = [30, 85, 14, 12, 27, 62, 89, 15, 78, 0]

sum_of_squares = 0

for x in x_list:
    # Square and add numbers
    sum_of_squares += delayed(my_square_function)(x)

result = sum_of_squares.compute()

# Print the answer
print(result)
```

27268

Sharing computation

```
delayed_intermediate = delayed(my_square_function)(3)

# These two results both use delayed_intermediate
delayed_result1 = delayed_intermediate - 5
delayed_result2 = delayed_intermediate + 4

# delayed_3_squared will be computed twice
print('delayed_result1:', delayed_result1.compute())
print('delayed_result2:', delayed_result2.compute())
```

```
delayed_result1: 4
delayed_result2: 13
```

Sharing computation

```
import dask

# delayed_intermediate will be computed once
comp_result1, comp_result2 = dask.compute(delayed_result1, delayed_result2)

print('comp_result1:', comp_result1)
print('comp_result2:', comp_result2)
```

```
delayed_result1: 4
delayed_result2: 13
```

Let's practice!

PARALLEL PROGRAMMING WITH DASK IN PYTHON

Task graphs and scheduling methods

PARALLEL PROGRAMMING WITH DASK IN PYTHON



James Fulton

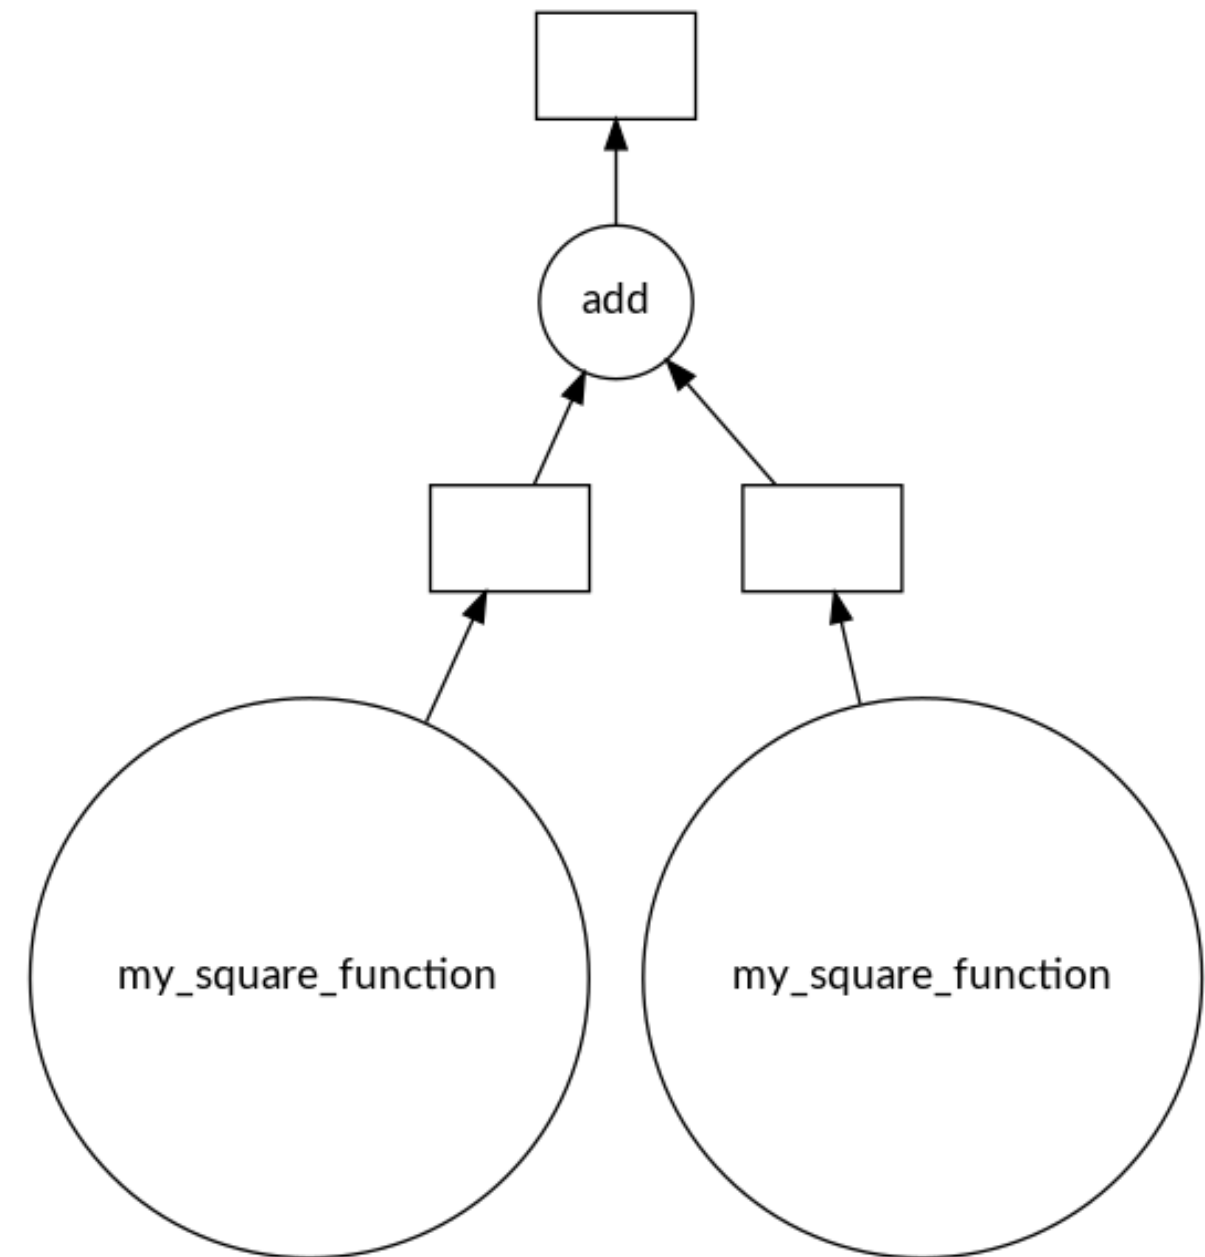
Climate Informatics Researcher

Visualizing a task graph

```
# Create 2 delayed objects
delayed_num1 = delayed(my_square_function)(3)
delayed_num2 = delayed(my_square_function)(4)

# Add them
result = delayed_num1 + delayed_num2

# Plot the task graph
result.visualize()
```




Overlapping task graph

```
delayed_intermediate = delayed(my_square_function)(3)


# These two results both use delayed_intermediate_result
delayed_result1 = delayed_intermediate - 5
delayed_result2 = delayed_intermediate + 4
```

Overlapping task graph

```
delayed_result1.visualize()
```




```
delayed_result2.visualize()
```



Overlapping task graph

```
# Plot the task graph  
dask.visualize(delayed_result1, delayed_result2)
```



Multi-threading vs. parallel processing

Moving data

Parallel processing

- Processes have their own RAM space

Multi-threading

- Threads use the same RAM space

Multi-threading vs. parallel processing

```
# Run a sum on two big arrays
```

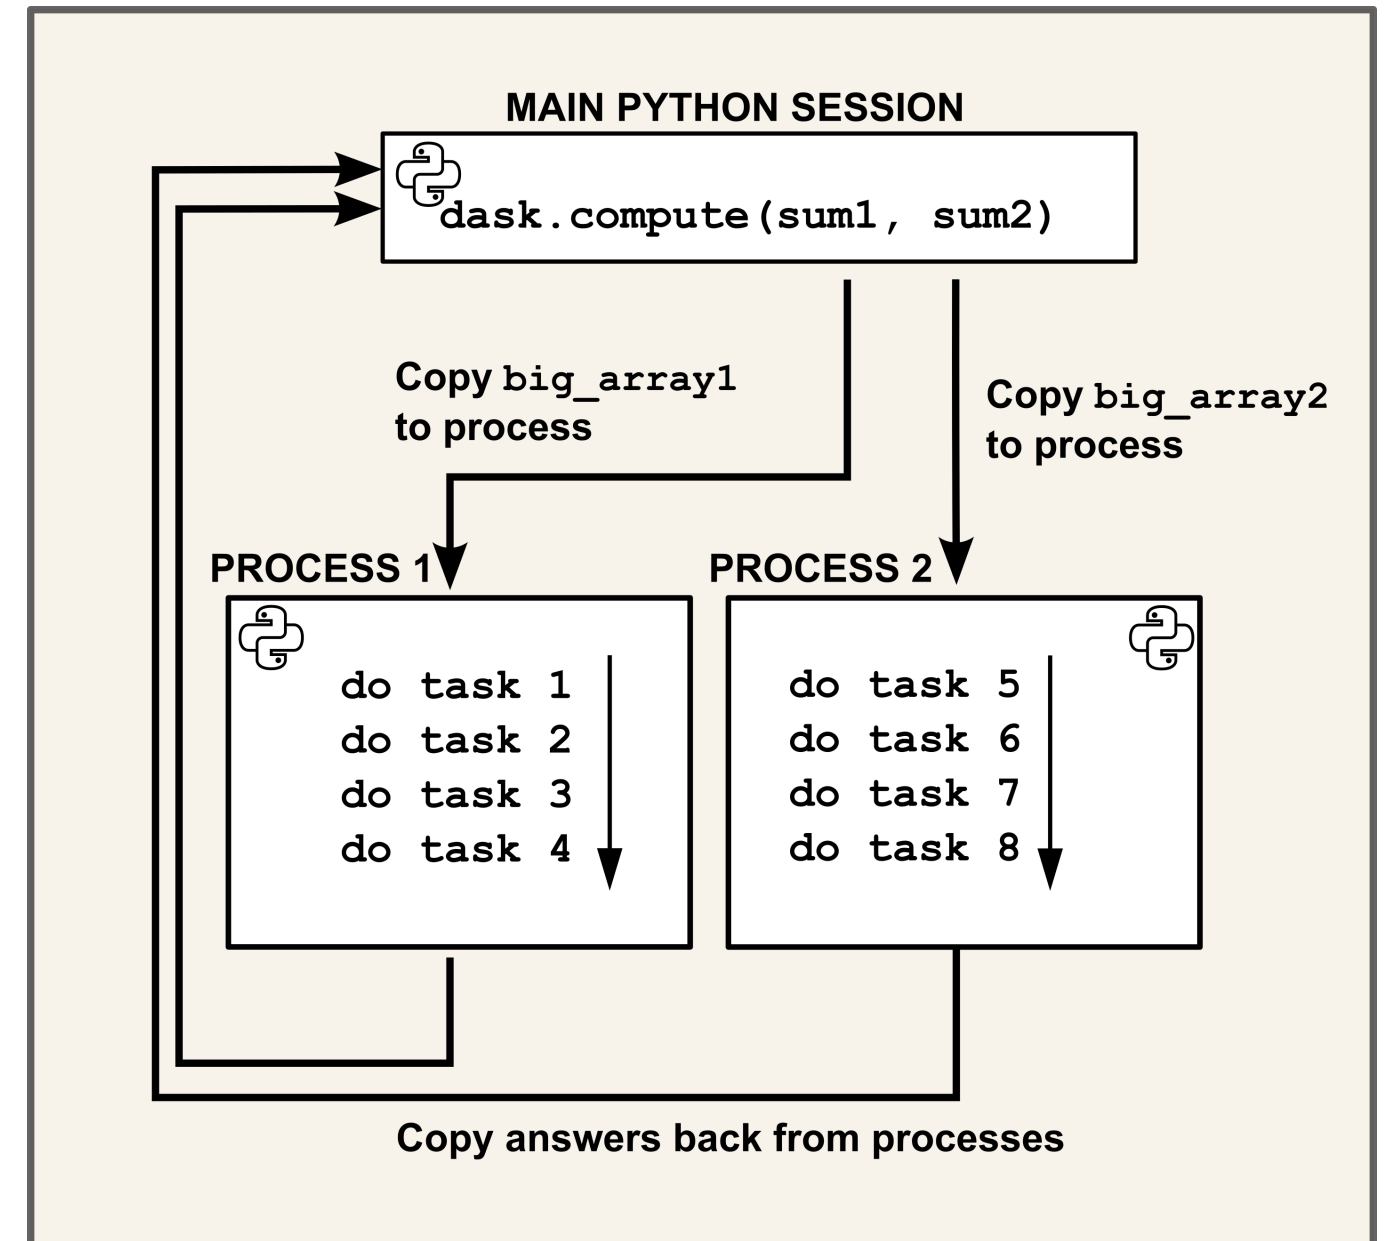
```
sum1 = delayed(np.sum)(big_array1)
```

```
sum2 = delayed(np.sum)(big_array2)
```

```
# Compute using processes
```

```
dask.compute(sum1, sum2)
```

- Slow using parallel processing



Multi-threading vs. parallel processing

```
# Run a sum on two big arrays
```

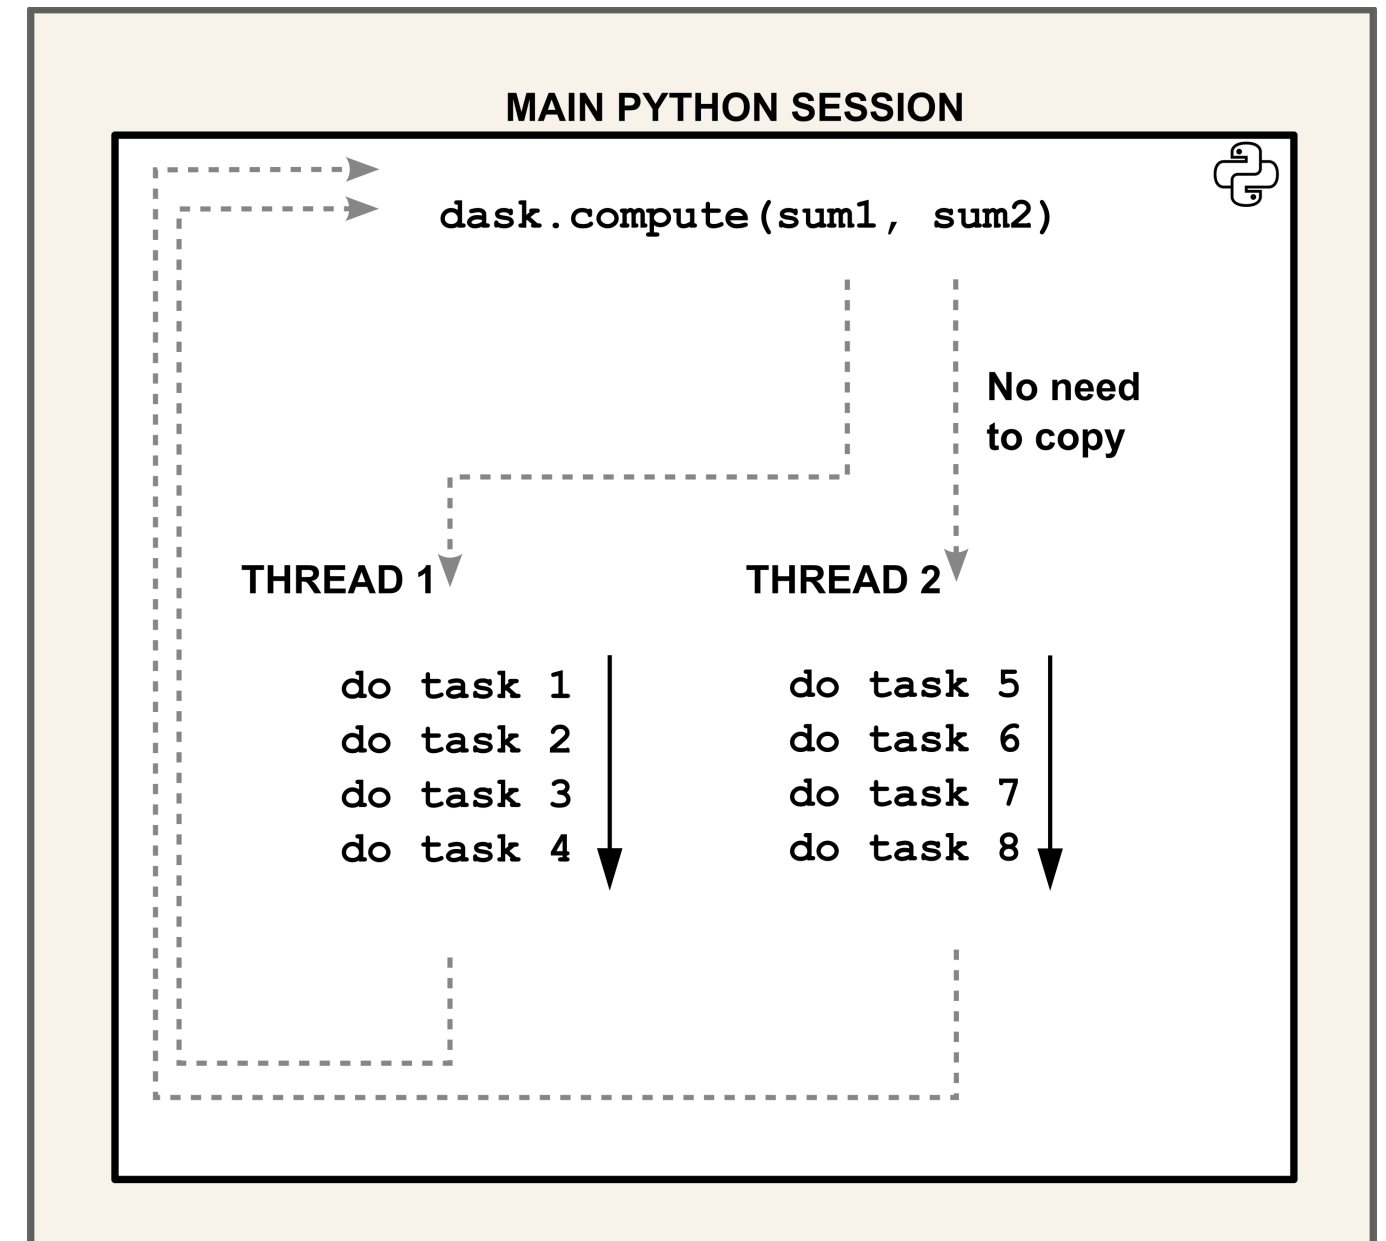
```
sum1 = delayed(np.sum)(big_array1)
```

```
sum2 = delayed(np.sum)(big_array2)
```

```
# Compute using threads
```

```
dask.compute(sum1, sum2)
```

- Fast using multi-threading



The GIL

Global interpreter lock - only one thread can read the Python script at a time



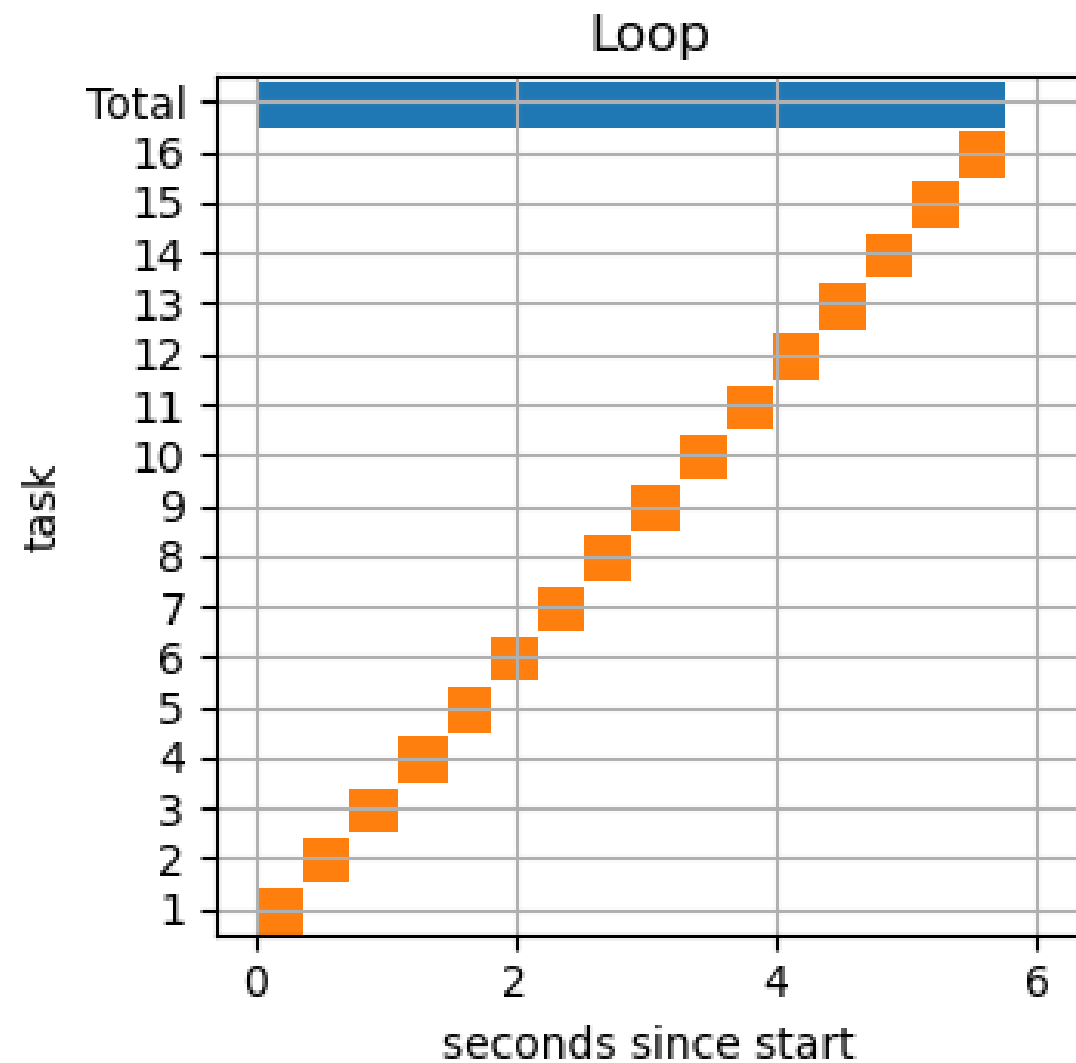
```
def sum_to_n(n):  
    """Sums numbers from 0 to n"""  
    total = 0  
    for i in range(n+1):  
        total += i  
    return total
```

- Multi-threading won't help here
- Parallel processing will

```
sum1 = delayed(sum_to_n)(1000)  
sum2 = delayed(sum_to_n)(1000)
```

Example timings - GIL

Basic Python loop



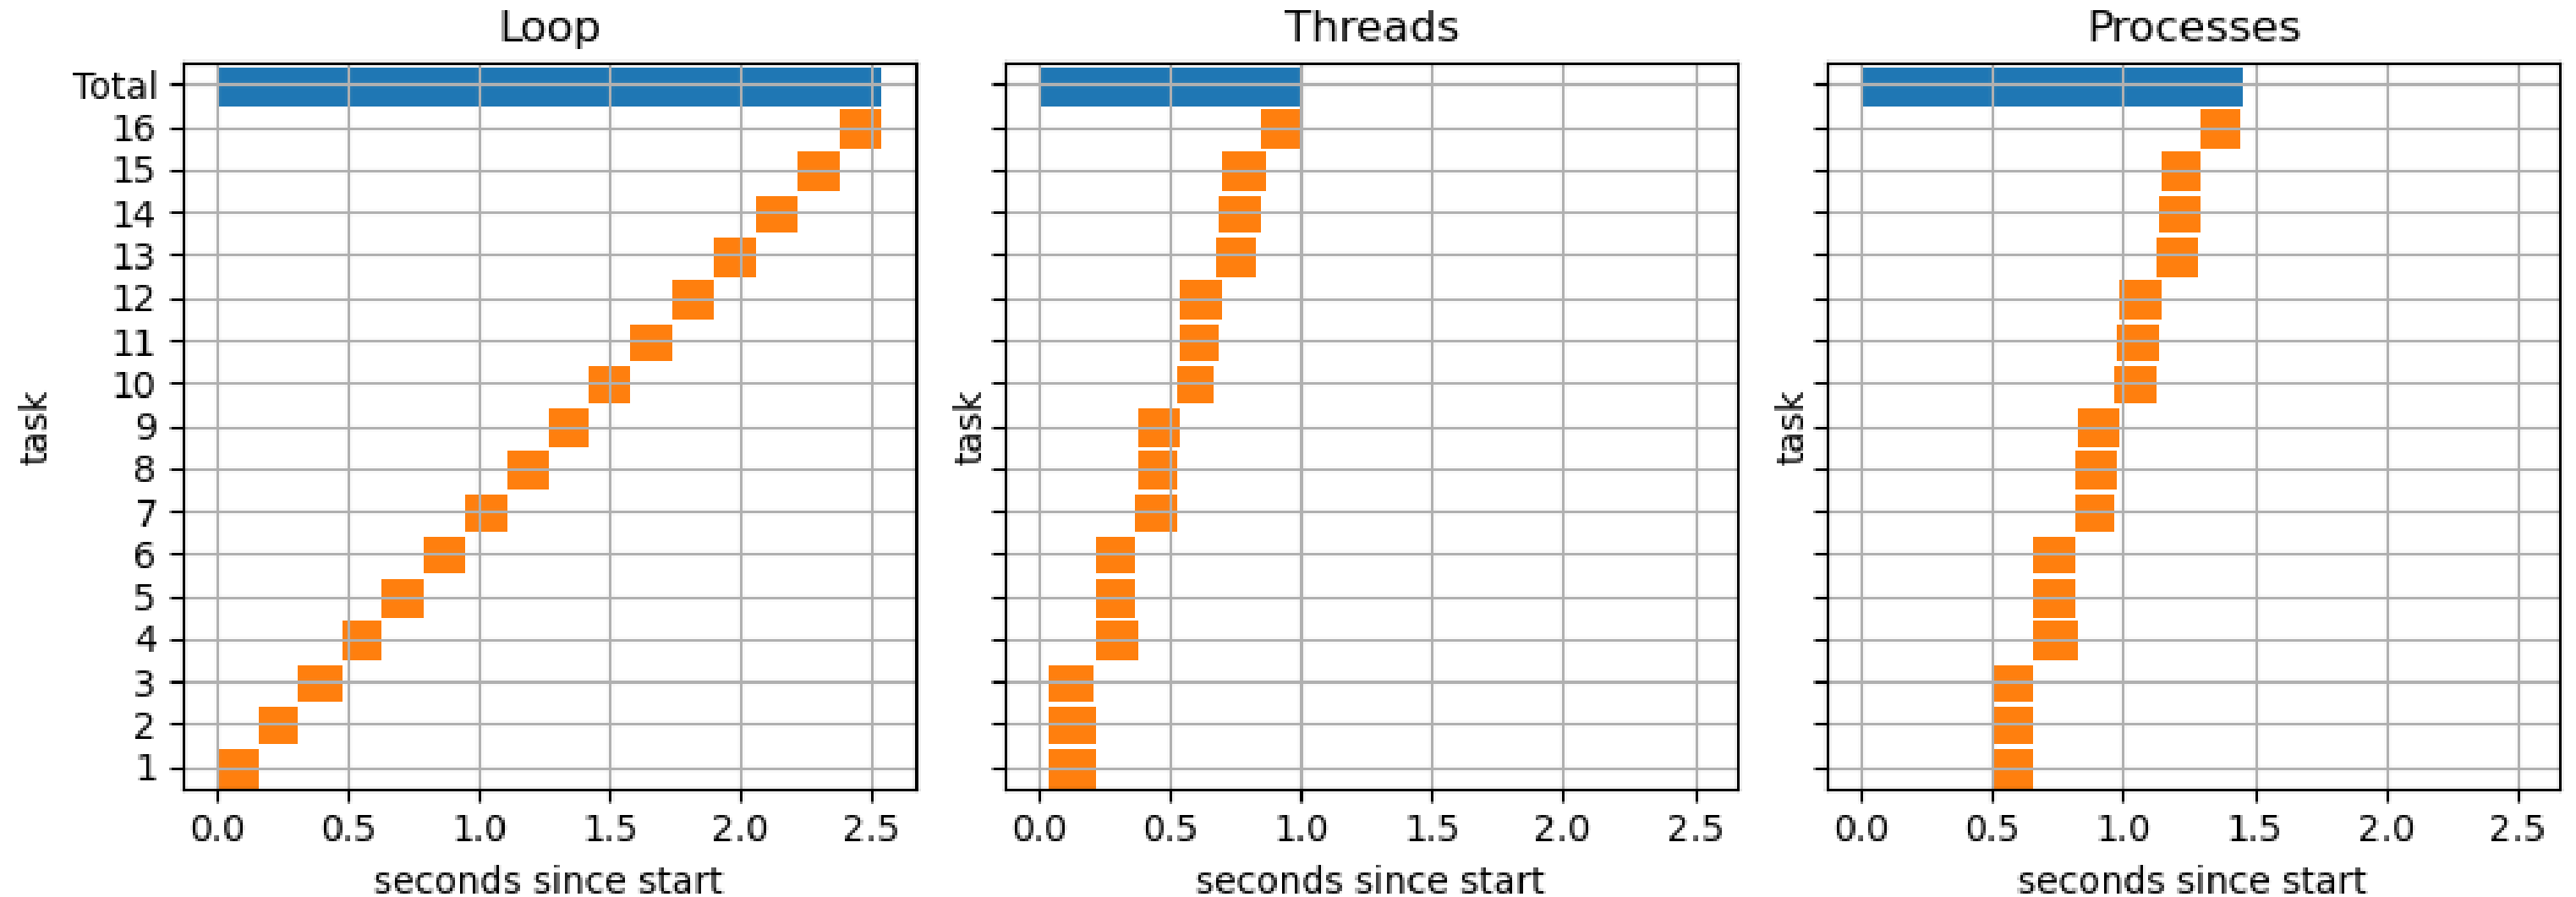
Functions which release the GIL

- E.g. the `pd.read_csv()` function releases the GIL

```
df1 = delayed(pd.read_csv)('file1.csv')  
df2 = delayed(pd.read_csv)('file2.csv')
```

Example timings - Loading data

Loading CSVs



Summary

Threads

- Are very fast to initiate
- Share memory space with main session
- No memory transfer needed
- Limited by the GIL, which allows one thread to read the code at once

Processes

- Take time and memory to set up
- Have separate memory pools
- Very slow to transfer data between themselves and to the main Python session
- Each have their own GIL and so don't need to take turns reading the code

Let's practice!

PARALLEL PROGRAMMING WITH DASK IN PYTHON

Building delayed pipelines

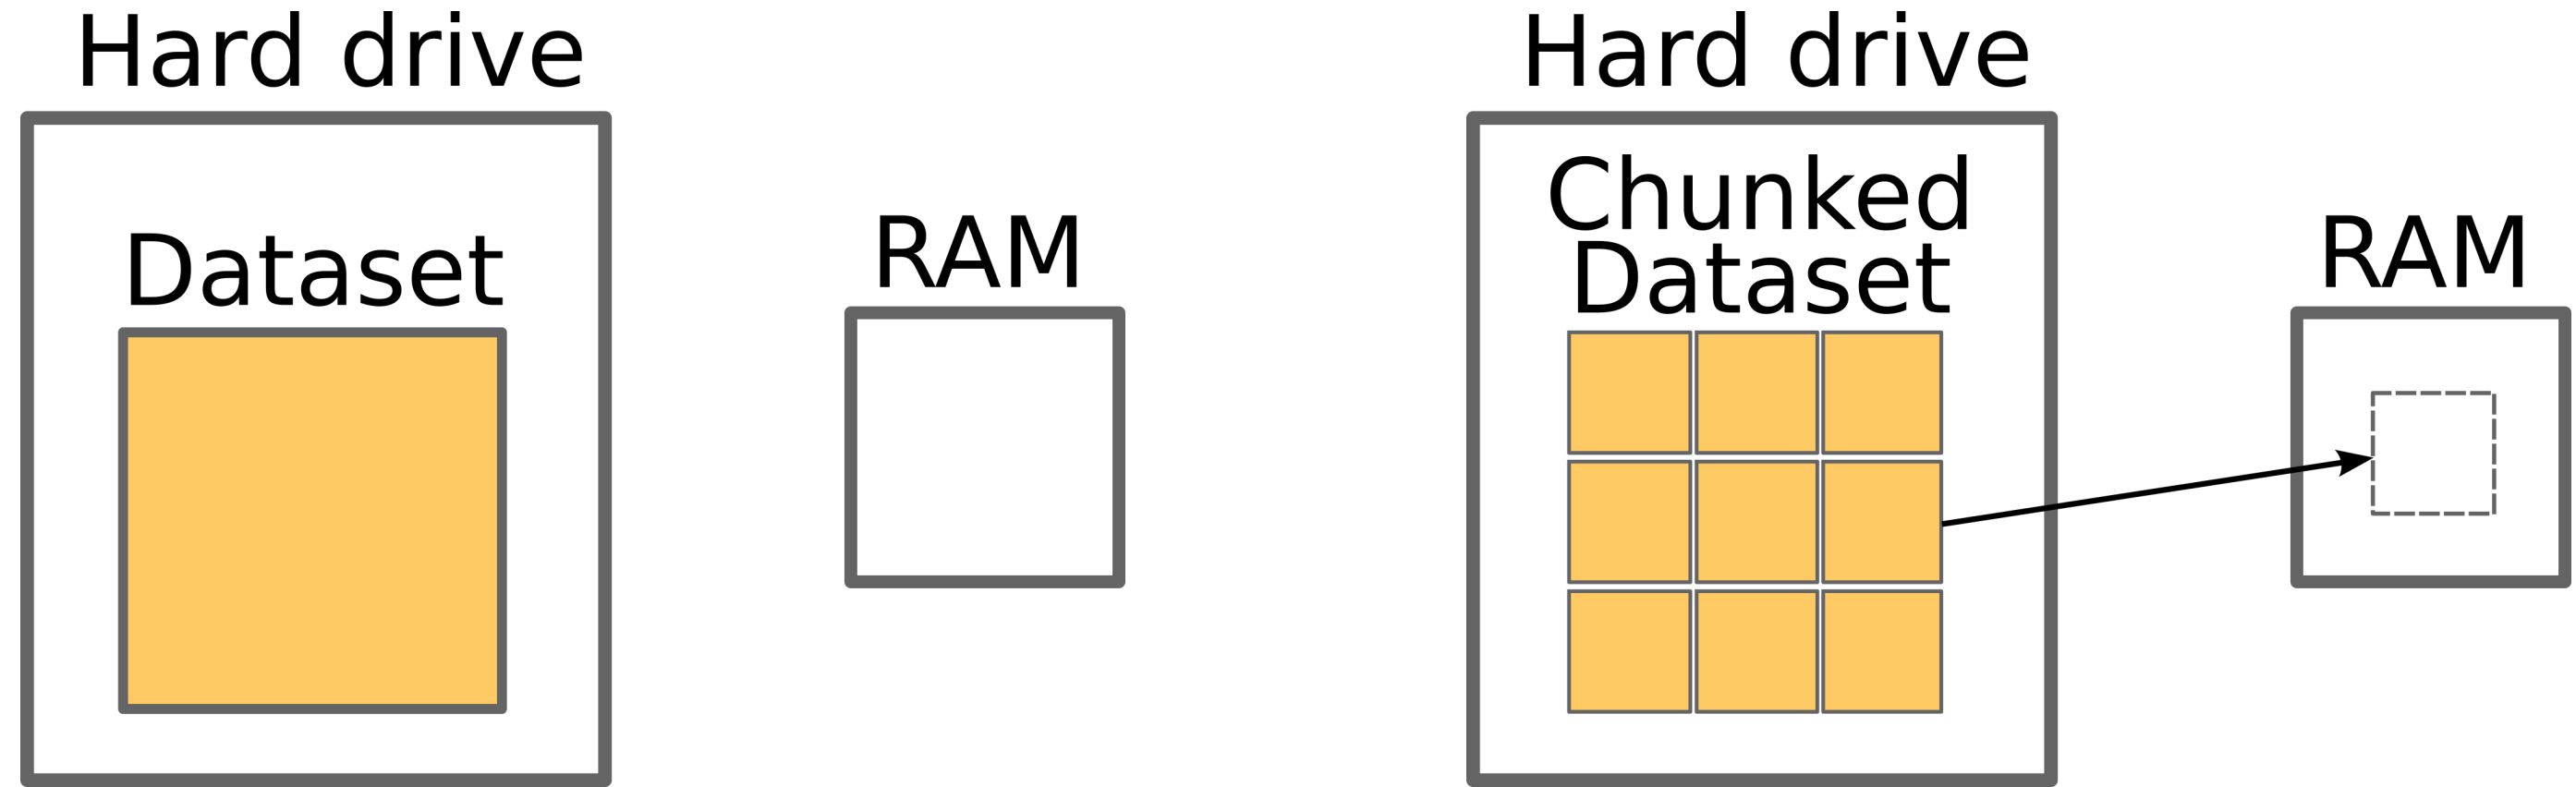
PARALLEL PROGRAMMING WITH DASK IN PYTHON



James Fulton

Climate Informatics Researcher

Chunks of data



Spotify songs dataset

```
files = [  
    '2005_tracks.csv',  
    '2006_tracks.csv',  
    '2007_tracks.csv',  
    '2008_tracks.csv',  
    '2009_tracks.csv',  
    '2010_tracks.csv',  
    ...  
    '2020_tracks.csv',  
]
```

Spotify songs dataset

```
      name  duration_ms  release_date  ...
0  Aldrig (feat. Carmon)    247869    2019-01-01  ...
2  2019 - The Year to Build    288105    2019-01-01  ...
3      Na zawsze    186812    2019-01-01  ...
4  Humo en la Trampa    258354    2019-01-01  ...
5      Au Au    176000    2019-01-01  ...
...      ...      ...      ...      ...
```

Analyzing the data

```
import pandas as pd

maximums = []

for file in files:
    # Load each file
    df = pd.read_csv(file)
    # Find maximum track length in each file
    max_length = df['duration_ms'].max()
    # Store this maximum
    maximums.append(max_length)

# Find the maximum of all the maximum lengths
absolute_maximum = max(maximums)
```

Analyzing the data

```
import pandas as pd

maximums = []

for file in files:
    # Load each file
    df = delayed(pd.read_csv)(file) # <----- delay loading
    # Find maximum track length in each file
    max_length = df['duration_ms'].max()
    # Store this maximum
    maximums.append(max_length)

# Find the maximum of all the maximum lengths
absolute_maximum = max(maximums)
```

Analyzing the data

```
import pandas as pd

maximums = []

for file in files:
    # Load each file
    df = delayed(pd.read_csv)(file) # <----- delay loading
    # Find maximum track length in each file
    max_length = df['duration_ms'].max()
    # Store this maximum
    maximums.append(max_length)

# Find the maximum of all the maximum lengths
absolute_maximum = delayed(max)(maximums) # <----- delay max() function
```


Using methods of a delayed object

```
import pandas as pd

maximums = []

for file in files:
    df = delayed(pd.read_csv)(file)
    # Use the .max() method
    max_length = df['duration_ms'].max()

    maximums.append(max_length)

absolute_maximum = delayed(max)(maximums)
```

```
print(max_length)
```

```
Delayed('max-0602855d-3ee6-4c43-a4d2-...')
```

- Delayed object methods and properties return new delayed objects

```
print(df.shape)
print(df.shape.compute())
```

```
Delayed('getattr-bc1e8838ab...')
(11907, 12)
```

Using methods of a delayed object

```
import pandas as pd

maximums = []

for file in files:
    df = delayed(pd.read_csv)(file)
    # Use a method which doesn't exist
    max_length = df['duration_ms'].fake()

    maximums.append(max_length)

absolute_maximum = delayed(max)(maximums)
```

```
print(max_length)
```

```
Delayed('max-6c026036-5daf-4b2-...')
```

- Methods aren't run until after `.compute()` is used

```
print(max_length.compute())
```

```
...
AttributeError: 'Series' object has no
attribute 'fake'
```

Using methods of a delayed object

```
import pandas as pd
```

```
maximums = []
```

```
for file in files:  
    df = delayed(pd.read_csv)(file)
```

```
    max_length = df['duration_ms'].max()  
    # Add delayed object to list  
    maximums.append(max_length)
```

```
# Run delayed max on delayed objects list  
absolute_maximum = delayed(max)(maximums)
```

`maximums` is a list of delayed objects

```
print(maximums)
```

```
[Delayed('max-80b...'),  
Delayed('max-fa15d...'),  
...]
```

Computing lists of delayed objects

```
import pandas as pd

maximums = []

for file in files:
    df = delayed(pd.read_csv)(file)

    max_length = df['duration_ms'].max()
    # Add delayed object to list
    maximums.append(max_length)

# Compute all the maximums
all_maximums = dask.compute(maximums)
```

```
print(all_maximums)
```

```
([2539418, 4368000, ...
... 4511716, 4864333],)
```

Computing lists of delayed objects

```
import pandas as pd

maximums = []

for file in files:
    df = delayed(pd.read_csv)(file)

    max_length = df['duration_ms'].max()

    maximums.append(max_length)

# Compute all the maximums
all_maximums = dask.compute(maximums)[0]
```

```
print(all_maximums)
```

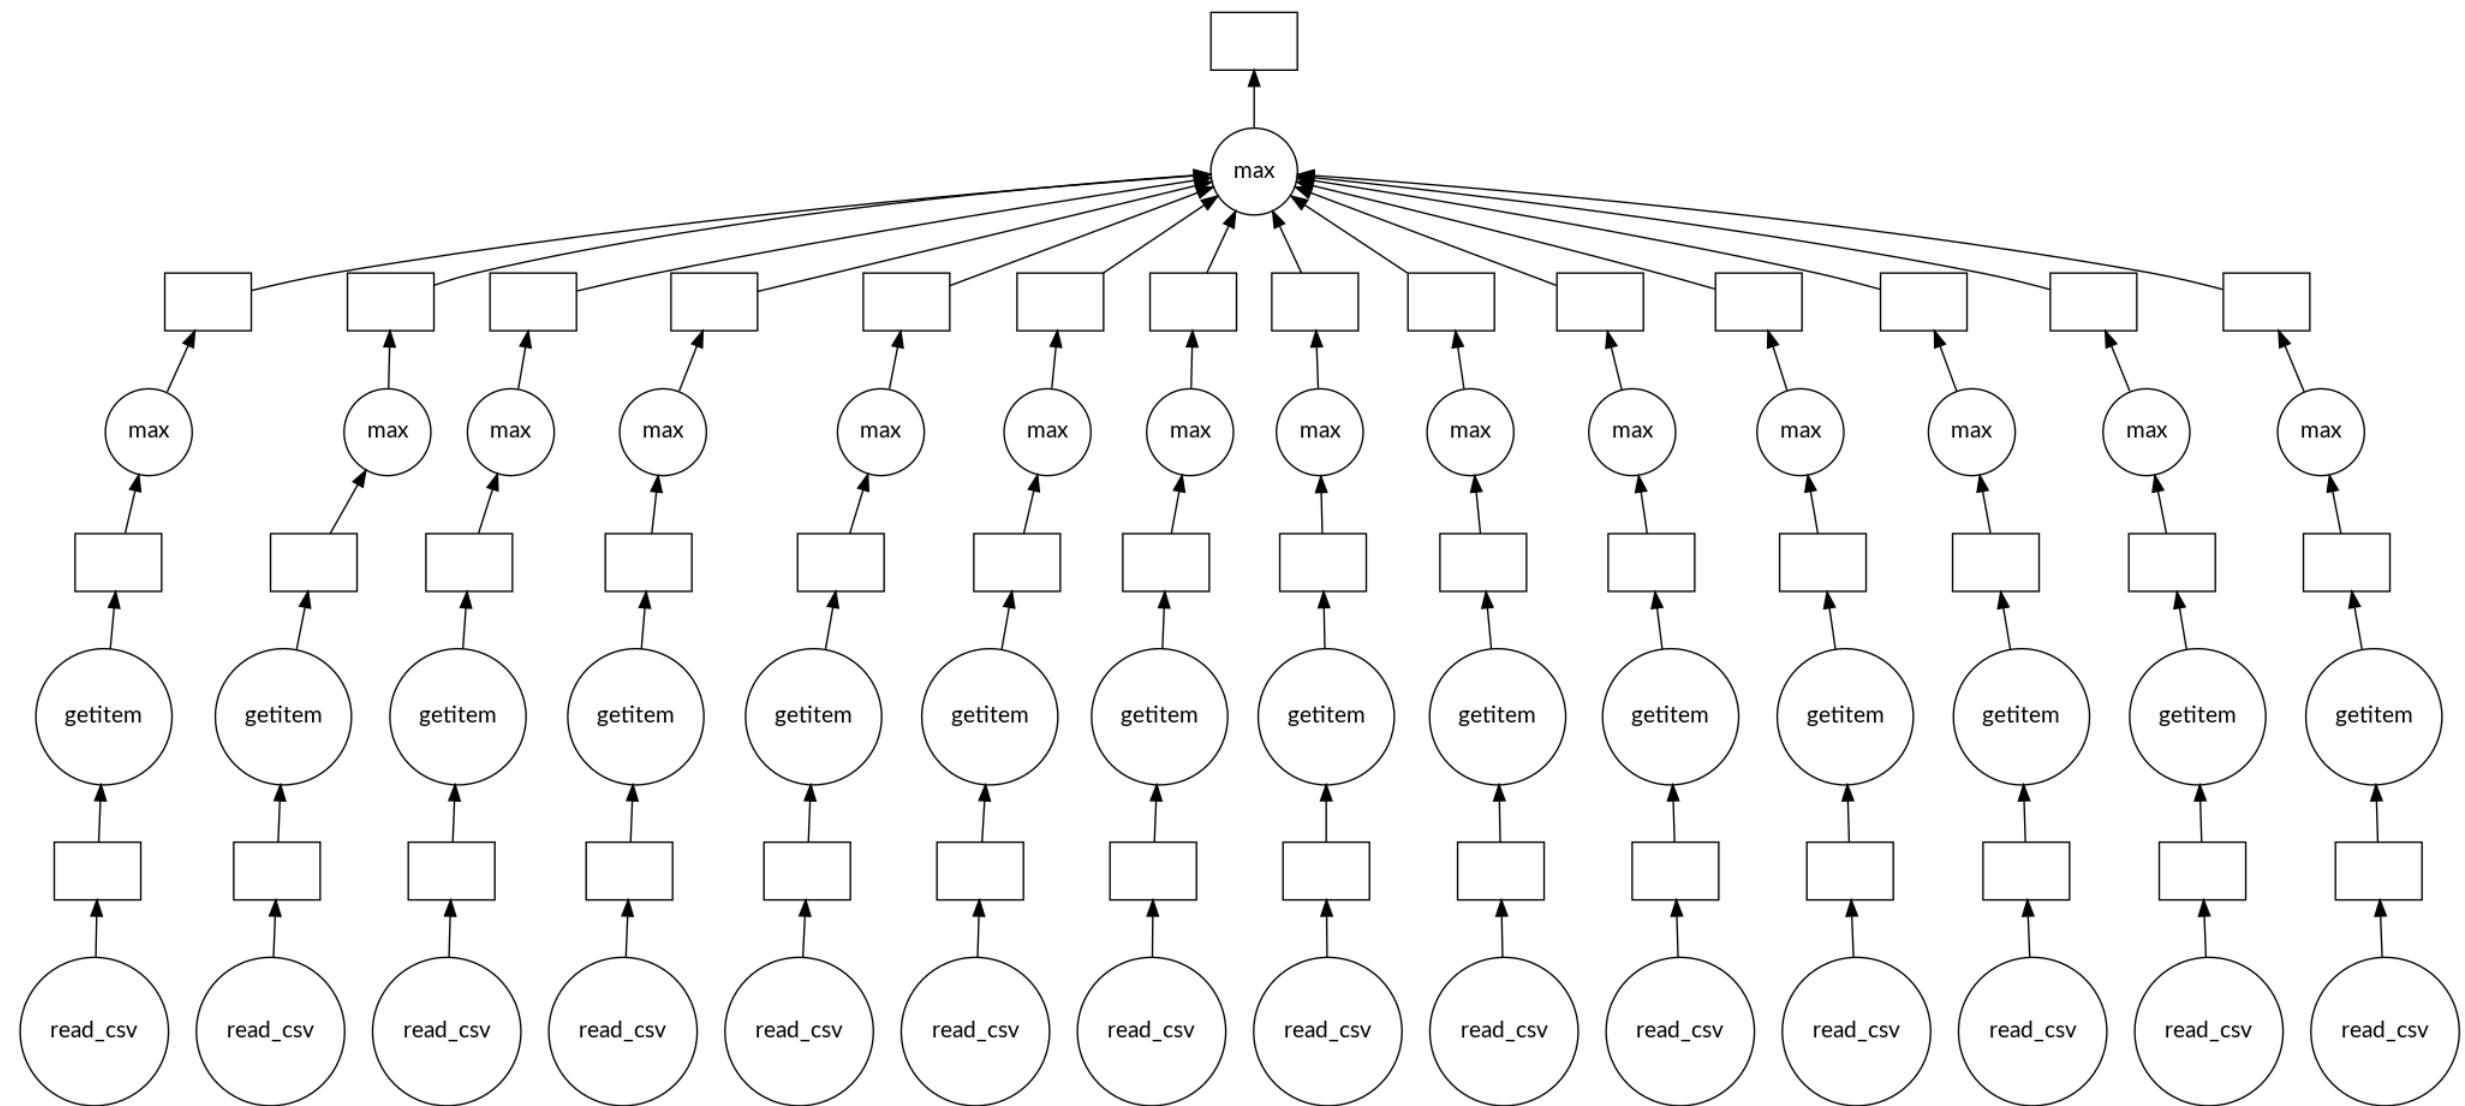
```
[2539418, 4368000, ...
... 4511716, 4864333]
```

To delay or not to delay

```
def get_max_track(df):  
    return df['duration_ms'].max()  
  
for file in files:  
    df = delayed(pd.read_csv)(file)  
    # Use function to find max  
    max_length = get_max_track(df)  
  
    maximums.append(max_length)  
  
absolute_maximum = delayed(max)(maximums)
```

Deeper task graph

```
absolute_maximum.visualize()
```



Let's practice!

PARALLEL PROGRAMMING WITH DASK IN PYTHON