Introduction to PySpark DataFrames

BIG DATA FUNDAMENTALS WITH PYSPARK

Upendra Devisetty
Science Analyst, CyVerse

What are PySpark DataFrames?

- PySpark SQL is a Spark library for structured data. It provides more information about the structure of data and computation

- PySpark DataFrame is an immutable distributed collection of data with named columns

- Designed for processing both structured (e.g relational database) and semi-structured data (e.g JSON)

- Dataframe API is available in Python, R, Scala, and Java

- DataFrames in PySpark support both SQL queries (`SELECT * from table`) or expression methods (`df.select()`)

SparkSession - Entry point for DataFrame API

- SparkContext is the main entry point for creating RDDs

- SparkSession provides a single point of entry to interact with Spark DataFrames

- SparkSession is used to create DataFrame, register DataFrames, execute SQL queries

- SparkSession is available in PySpark shell as `spark`

Creating DataFrames in PySpark

- Two different methods of creating DataFrames in PySpark
  - From existing RDDs using SparkSession's createDataFrame() method

  - From various data sources (CSV, JSON, TXT) using SparkSession's read method

- Schema controls the data and helps DataFrames to optimize queries

- Schema provides information about column name, type of data in the column, empty values etc.,

Create a DataFrame from RDD

```python
iphones_RDD = sc.parallelize([
    ("XS", 2018, 5.65, 2.79, 6.24),
    ("XR", 2018, 5.94, 2.98, 6.84),
    ("X10", 2017, 5.65, 2.79, 6.13),
    ("8Plus", 2017, 6.23, 3.07, 7.12)
])
```

```python
names = ['Model', 'Year', 'Height', 'Width', 'Weight']
```

```python
iphones_df = spark.createDataFrame(iphones_RDD, schema=names)
type(iphones_df)
```

```
pyspark.sql.dataframe.DataFrame
```

Create a DataFrame from reading a CSV/JSON/TXT

```
df_csv = spark.read.csv("people.csv", header=True, inferSchema=True)
```

```
df_json = spark.read.json("people.json")
```

```
df_txt = spark.read.txt("people.txt")
```

- Path to the file and two optional parameters

- Two optional parameters
  - `header=True` , `inferSchema=True`

Let's practice

BIG DATA FUNDAMENTALS WITH PYSPARK

Interacting with PySpark DataFrames

BIG DATA FUNDAMENTALS WITH PYSPARK

Upendra Devisetty
Science Analyst, CyVerse

DataFrame operators in PySpark

- DataFrame operations: Transformations and Actions

- DataFrame Transformations:
  - select(), filter(), groupby(), orderby(), dropDuplicates(), and withColumnRenamed()

- DataFrame Actions :
  - printSchema(), head(), show(), count(), columns, and describe()

  Correction: printSchema() is a method for any Spark dataset/dataframe and not an action

select() and show() operations

- `select()` transformation subsets the columns in the DataFrame

```
df_id_age = test.select('Age')
```

- `show()` action prints first 20 rows in the DataFrame

```
df_id_age.show(3)
```

```
+---+
|Age|
+---+
| 17|
| 17|
| 17|
+---+
only showing top 3 rows
```

filter() and show() operations

- `filter()` transformation filters out the rows based on a condition

```
new_df_age21 = new_df.filter(new_df.Age > 21)
new_df_age21.show(3)
```

```
+-------+------+---+
|User_ID|Gender|Age|
+-------+------+---+
|1000002|     M| 55|
|1000003|     M| 26|
|1000004|     M| 46|
+-------+------+---+
only showing top 3 rows
```

groupby() and count() operations

- `groupby()` operation can be used to group a variable

```
test_df_age_group = test_df.groupby('Age')
test_df_age_group.count().show(3)
```

```
+---+------+
|Age| count|
+---+------+
| 26|219587|
| 17|     4|
| 55| 21504|
+---+------+
only showing top 3 rows
```

orderby() Transformations

- `orderby()` operation sorts the DataFrame based on one or more columns

```
test_df_age_group.count().orderBy('Age').show(3)
```

```
+---+-----+
|Age|count|
+---+-----+
|  0|15098|
| 17|    4|
| 18|99660|
+---+-----+
only showing top 3 rows
```

dropDuplicates()

- `dropDuplicates()` removes the duplicate rows of a DataFrame

```
test_df_no_dup = test_df.select('User_ID','Gender', 'Age').dropDuplicates()
test_df_no_dup.count()
```

```
5892
```

withColumnRenamed Transformations

- `withColumnRenamed()` renames a column in the DataFrame

```python
test_df_sex = test_df.withColumnRenamed('Gender', 'Sex')
test_df_sex.show(3)
```

```
+-------+---+---+

|User_ID|Sex|Age|

+-------+---+---+

|1000001|  F| 17|
|1000001|  F| 17|
|1000001|  F| 17|

+-------+---+---+
```

printSchema()

- `printSchema()` operation prints the types of columns in the DataFrame

```
test_df.printSchema()
```

```
|-- User_ID: integer (nullable = true)
|-- Product_ID: string (nullable = true)
|-- Gender: string (nullable = true)
|-- Age: string (nullable = true)
|-- Occupation: integer (nullable = true)
|-- Purchase: integer (nullable = true)
```

columns actions

- `columns` operator prints the columns of a DataFrame

```
test_df.columns
```

```
['User_ID', 'Gender', 'Age']
```

describe() actions

- `describe()` operation compute summary statistics of numerical columns in the DataFrame

```
test_df.describe().show()
```

```
+-------+------------------+------+------------------+
|summary|           User_ID|Gender|               Age|
+-------+------------------+------+------------------+
|  count|            550068|550068|            550068|
|   mean|1003028.8424013031|  null|30.382052764385495|
| stddev|1727.5915855307312|  null|11.866105189533554|
|    min|           1000001|     F|                 0|
|    max|           1006040|     M|                55|
+-------+------------------+------+------------------+
```

Let's practice

BIG DATA FUNDAMENTALS WITH PYSPARK

DataFrame API vs SQL queries

- In PySpark You can interact with SparkSQL through DataFrame API and SQL queries

- The DataFrame API provides a programmatic domain-specific language (DSL) for data

- DataFrame transformations and actions are easier to construct programmatically

- SQL queries can be concise and easier to understand and portable

- The operations on DataFrames can also be done using SQL queries

Executing SQL Queries

- The SparkSession `sql()` method executes SQL query

- `sql()` method takes a SQL statement as an argument and returns the result as DataFrame

```
df.createOrReplaceTempView("table1")
```

```
df2 = spark.sql("SELECT field1, field2 FROM table1")
df2.collect()
```

```
[Row(f1=1, f2='row1'), Row(f1=2, f2='row2'), Row(f1=3, f2='row3')]
```

SQL query to extract data

```
test_df.createOrReplaceTempView("test_table")
```

```
query = '''SELECT Product_ID FROM test_table'''
```

```
test_product_df = spark.sql(query)
test_product_df.show(5)
```

```
+---------+
|Product_ID|
+---------+
| P00069042|
| P00248942|
| P00087842|
| P00085442|
| P00285442|
+---------+
```

Summarizing and grouping data using SQL queries

```
test_df.createOrReplaceTempView("test_table")
```

```
query = '''SELECT Age, max(Purchase) FROM test_table GROUP BY Age'''
```

```
spark.sql(query).show(5)
```

```
+-----+------------+
|  Age|max(Purchase)|
+-----+------------+
|18-25|       23958|
|26-35|       23961|
| 0-17|       23955|
|46-50|       23960|
|51-55|       23960|
+-----+------------+
only showing top 5 rows
```

Filtering columns using SQL queries

```
test_df.createOrReplaceTempView("test_table")
```

```
query = '''SELECT Age, Purchase, Gender FROM test_table WHERE Purchase > 20000 AND Gender == "F"'''
```

```
spark.sql(query).show(5)
```

```
+-----+--------+------+
|  Age|Purchase|Gender|
+-----+--------+------+
|36-45|   23792|     F|
|26-35|   21002|     F|
|26-35|   23595|     F|
|26-35|   23341|     F|
|46-50|   20771|     F|
+-----+--------+------+
only showing top 5 rows
```
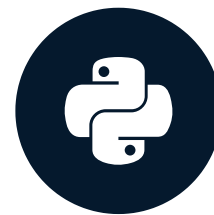
Time to practice!

datacamp

What is Data visualization?

- Data visualization is a way of representing your data in graphs or charts

- Open source plotting tools to aid visualization in Python:
  - Matplotlib, Seaborn, Bokeh etc.,

- Plotting graphs using PySpark DataFrames is done using three methods
  - pyspark_dist_explore library

  - toPandas()

  - HandySpark library

Data Visualization using Pyspark_dist_explore

- `Pyspark_dist_explore` library provides quick insights into DataFrames

- Currently three functions available : `hist()` , `distplot()` , and `pandas_histogram()`

```python
test_df = spark.read.csv("test.csv", header=True, inferSchema=True)
```

```python
test_df_age = test_df.select('Age')
```

```python
hist(test_df_age, bins=20, color="red")
```

Using Pandas for plotting DataFrames

- It's easy to create charts from pandas DataFrames

```
test_df = spark.read.csv("test.csv", header=True, inferSchema=True)
```

```
test_df_sample_pandas = test_df.toPandas()
```

```
test_df_sample_pandas.hist('Age')
```

- **Note:** When you have large volumes of data, using `toPandas()` isn't recommended

Pandas DataFrame vs PySpark DataFrame

- Pandas DataFrames are in-memory, single-server based structures and operations on PySpark run in parallel

- The result is generated as we apply any operation in Pandas whereas operations in PySpark DataFrame are lazy evaluation

- Pandas DataFrame as mutable and PySpark DataFrames are immutable

- Pandas API support more operations than PySpark Dataframe API

HandySpark method of visualization

- HandySpark is a package designed to improve PySpark user experience
    - Easy data fetching

    - Distributed computation retained

```python
test_df = spark.read.csv('test.csv', header=True, inferSchema=True)
```

```python
hdf = test_df.toHandy()
```

```python
hdf.cols["Age"].hist()
```

Let's visualize DataFrames

datacamp