

What is a column family database?

NOSQL CONCEPTS




Miriam Antona
Software engineer

Column family databases - overview

- Derive from **Google BigTable**
- Store data in **column families**
 - group related data
 - frequently accessed together
- Also called **wide column** databases
- Great when dealing with large volumes of data

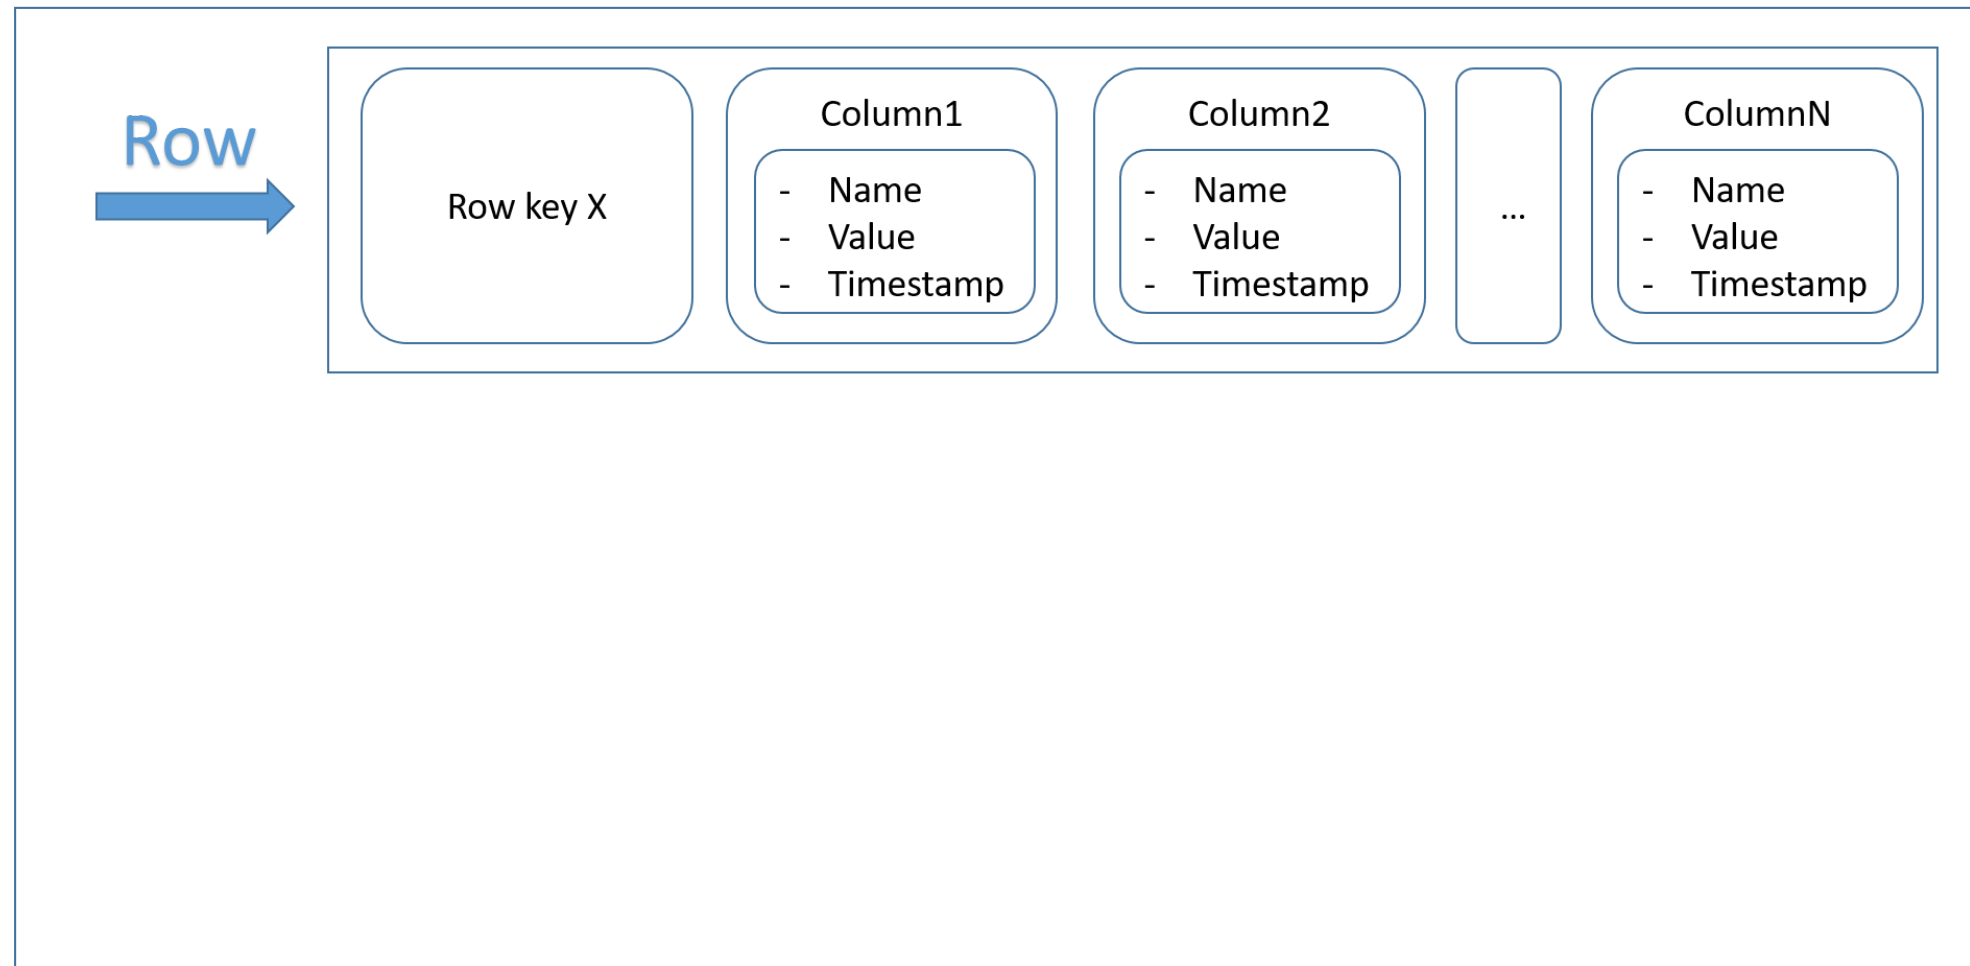
Column family databases - structure

Column family



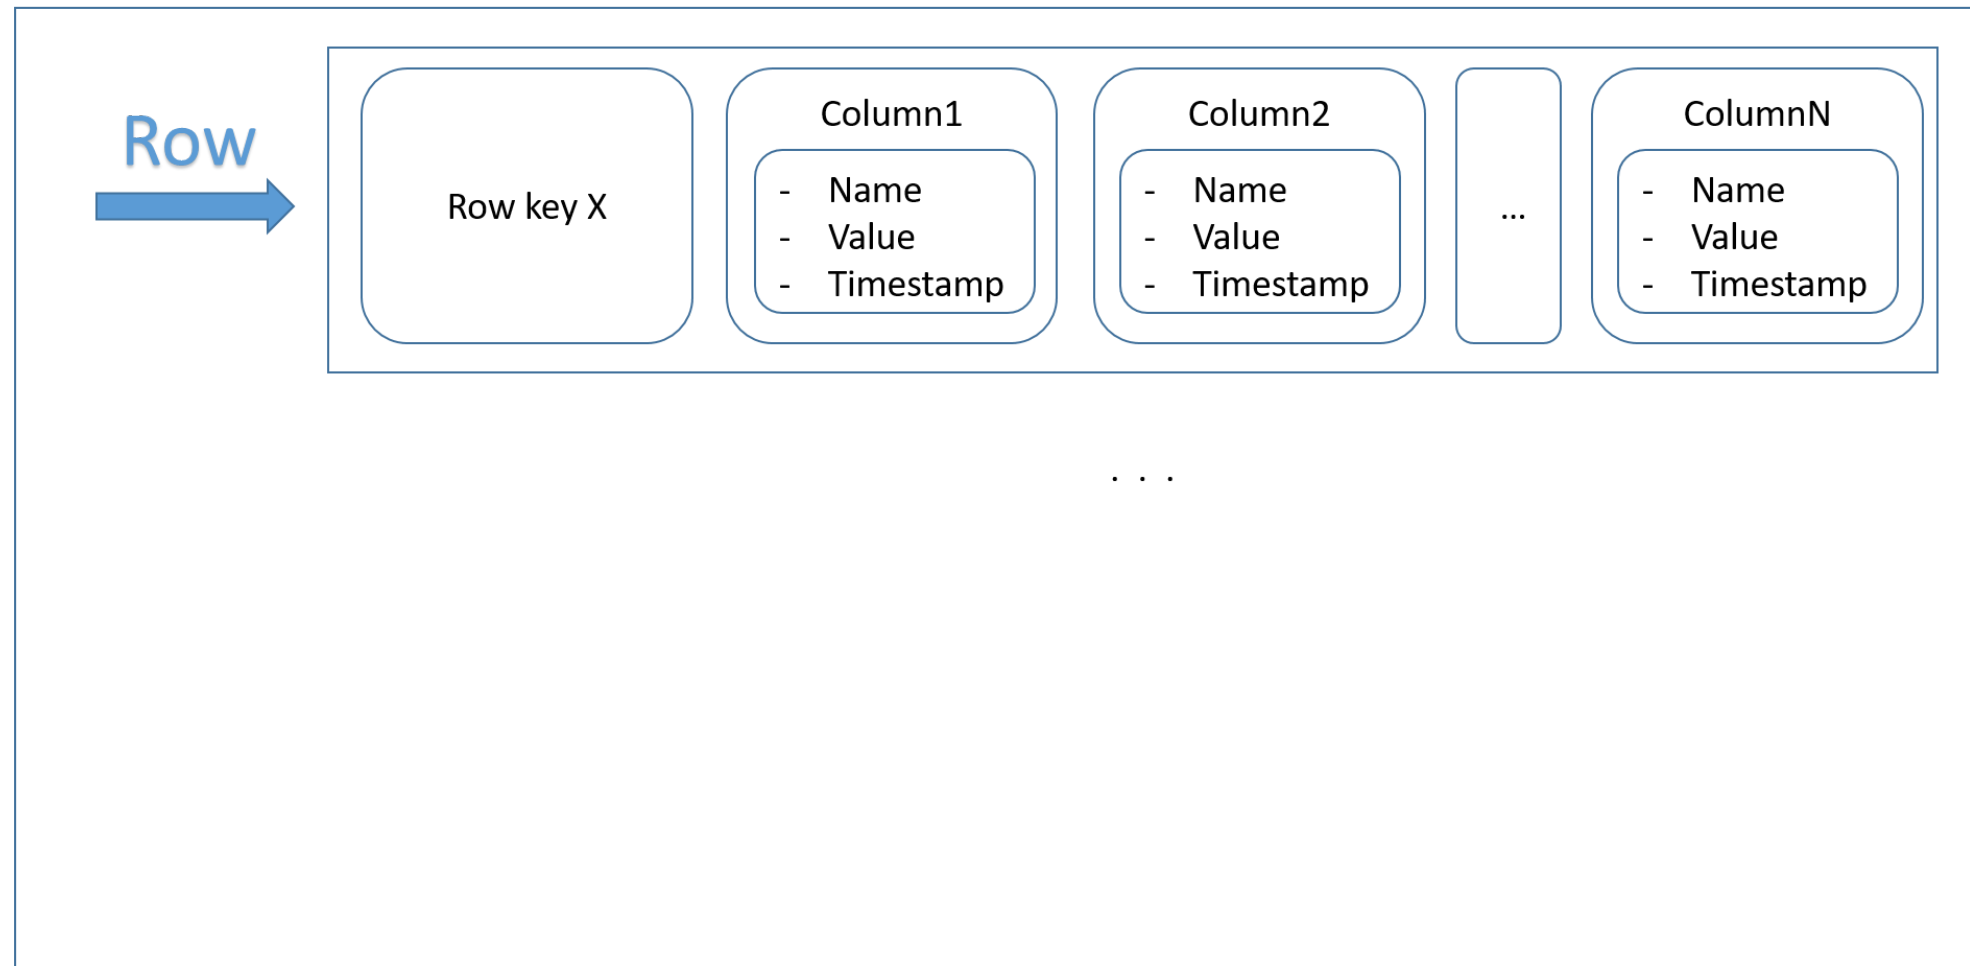
Column family databases - structure

Column family



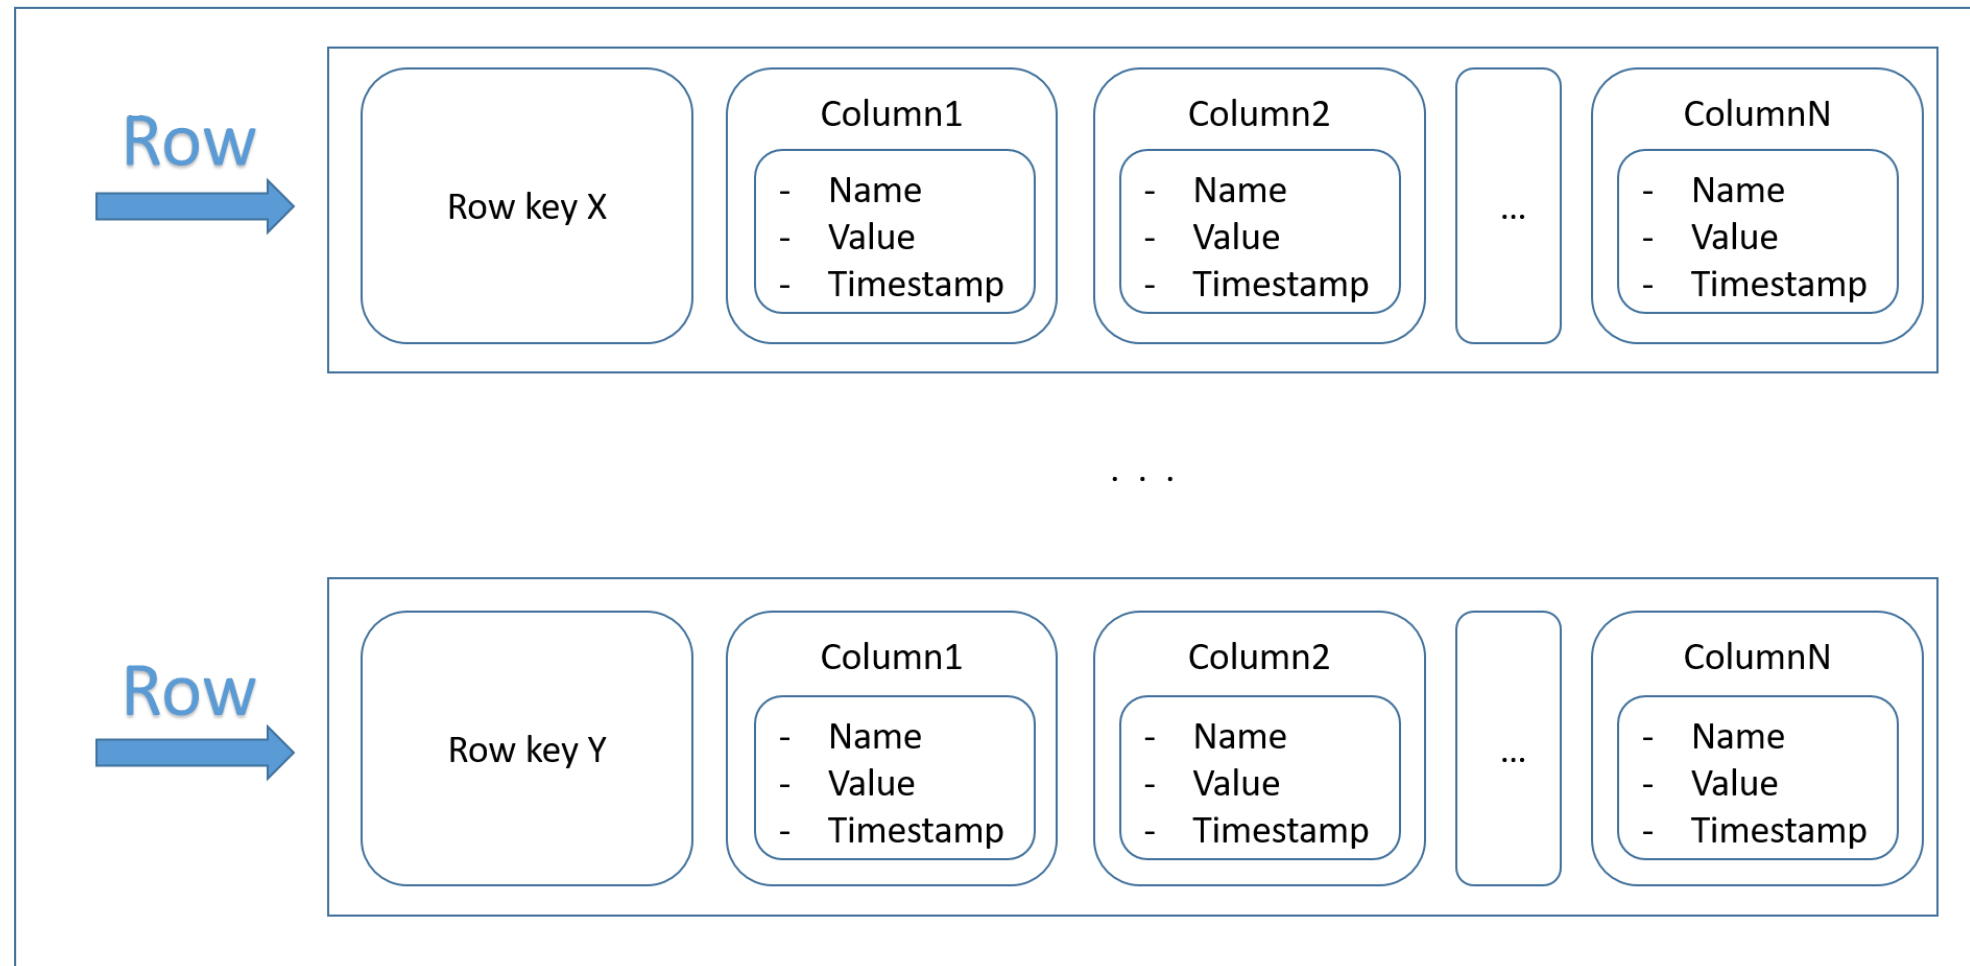
Column family databases - structure

Column family



Column family databases - structure

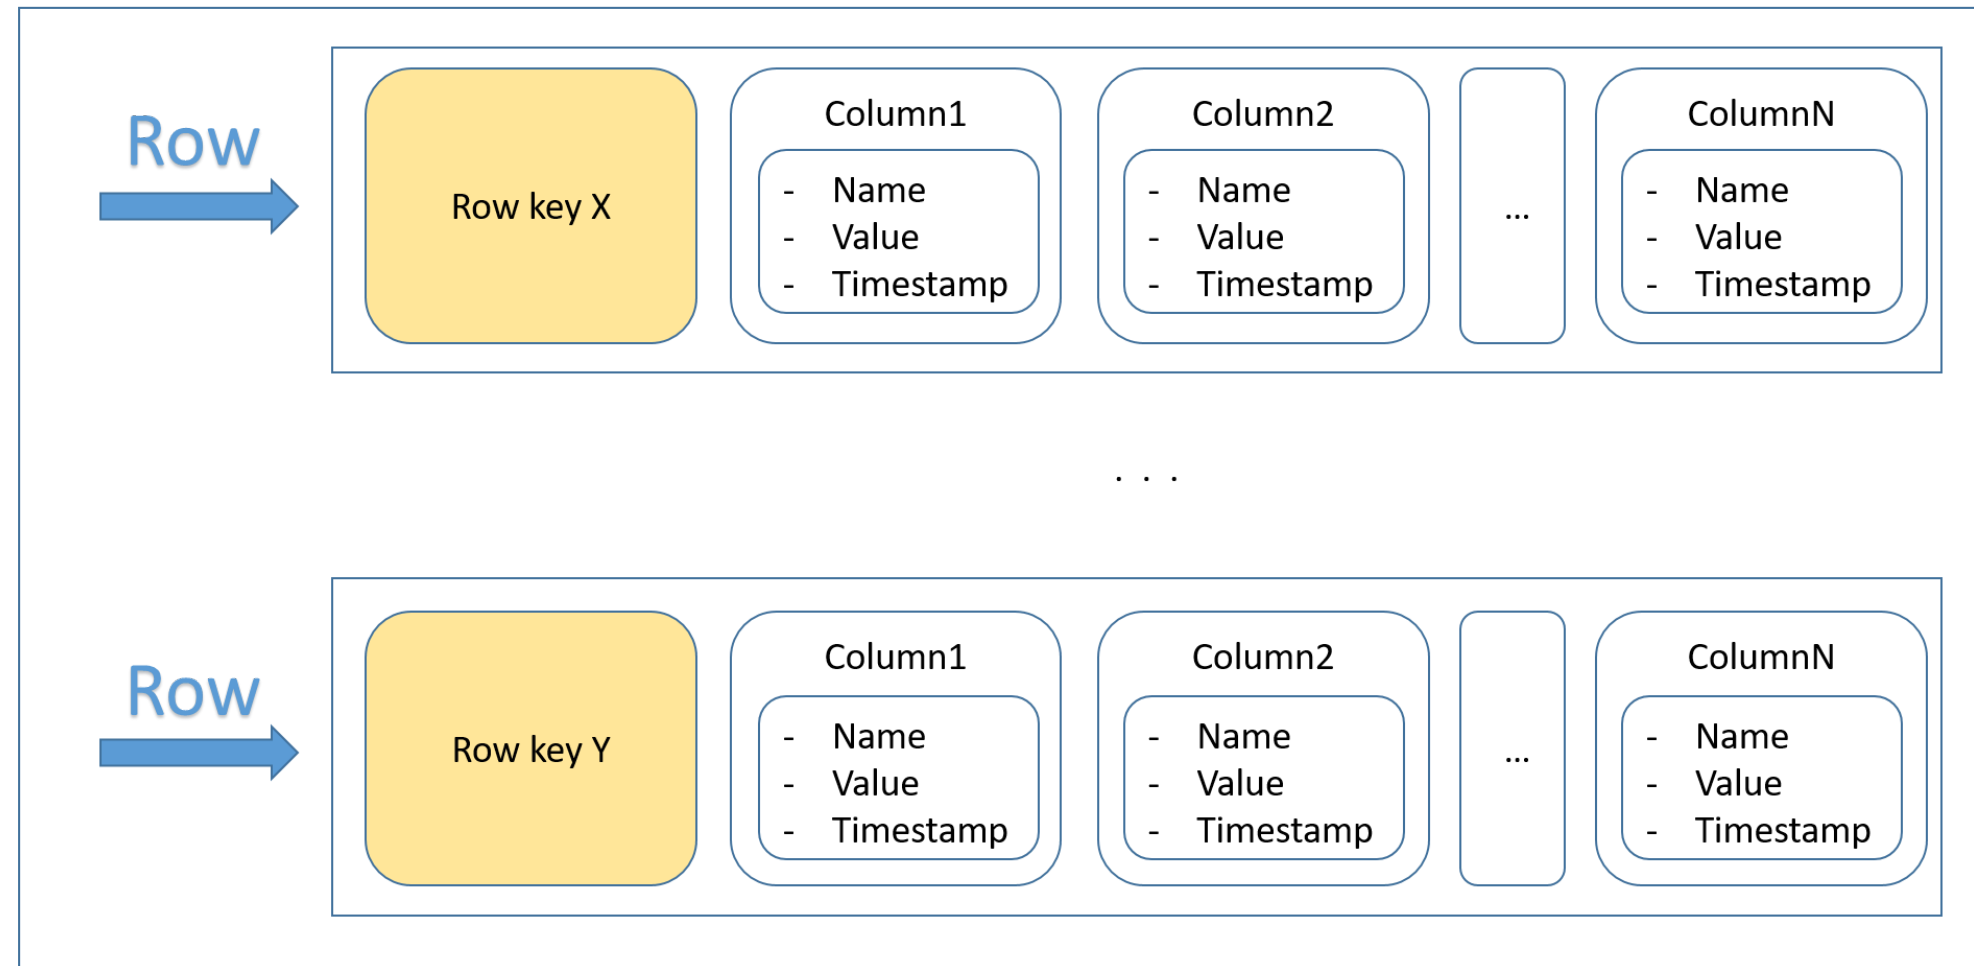
Column family



- A **column family** is like a **table** in a relational database

Column family databases - structure

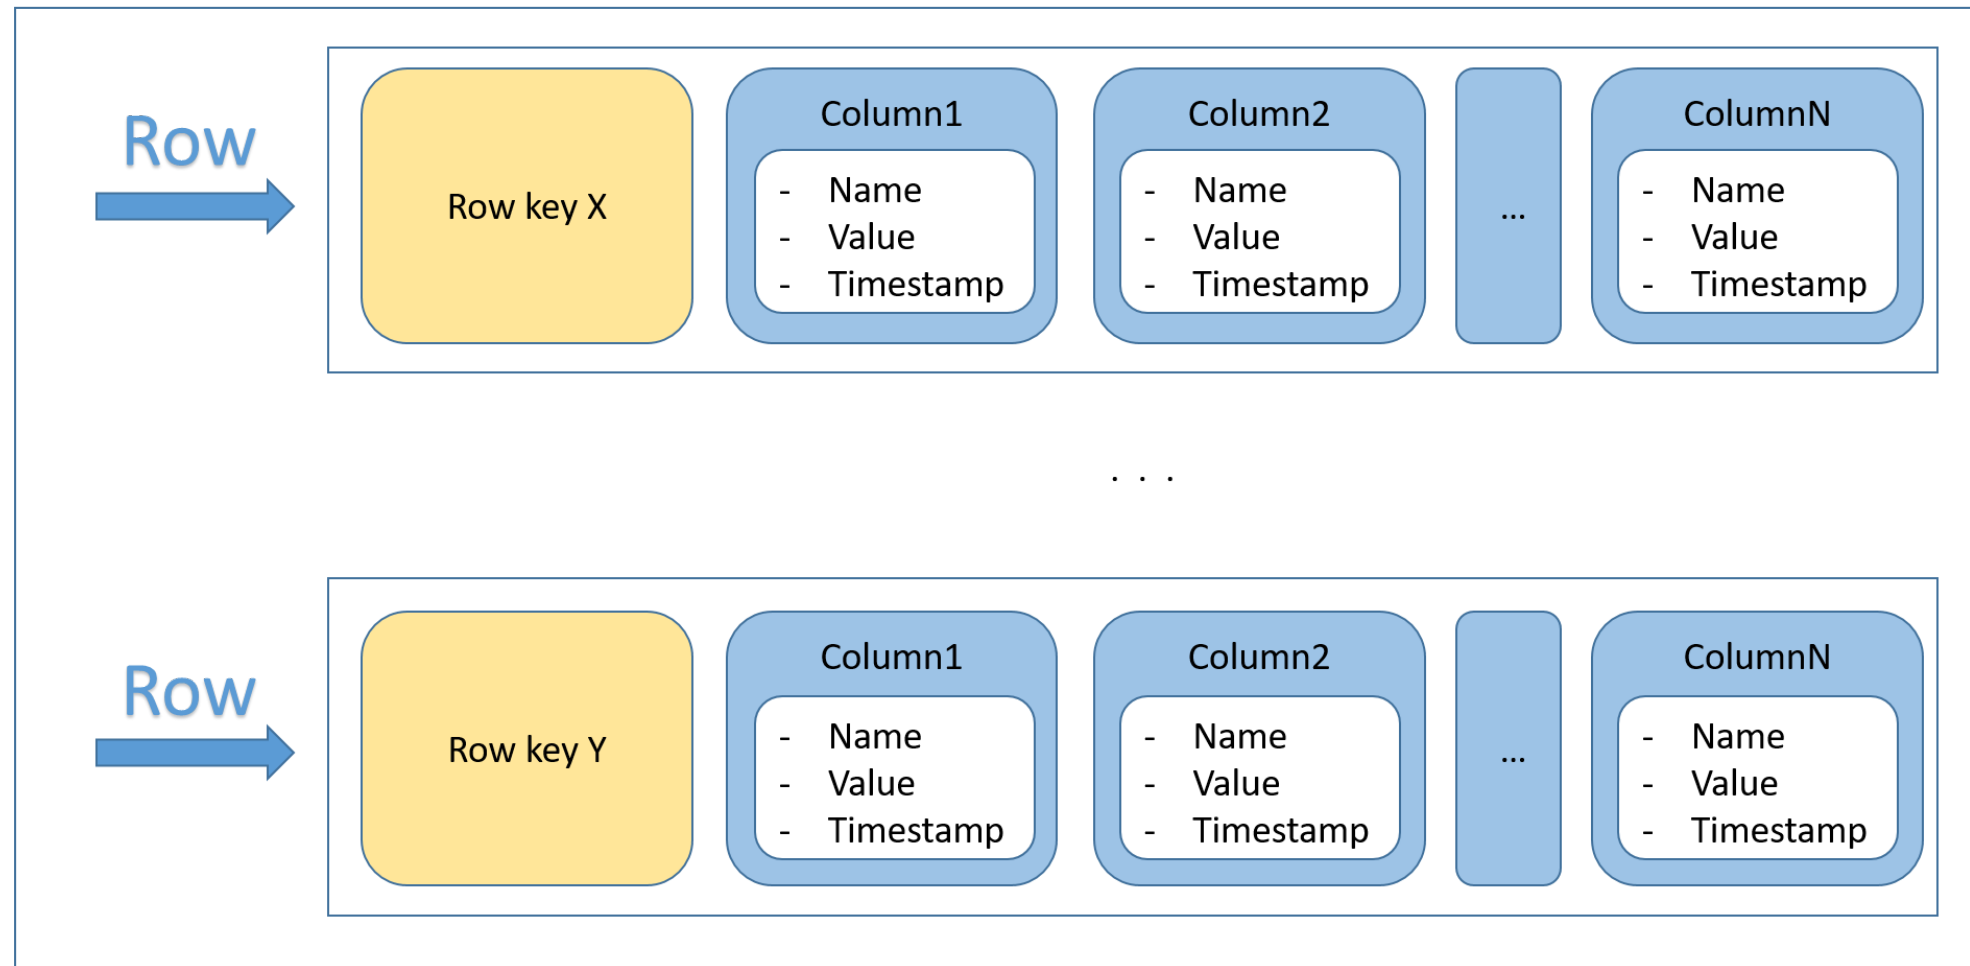
Column family



- **Row key:** unique identifiers
 - Like primary keys in a relational database

Column family databases - structure


Column family



- Each **row** can have different number of columns
 - Columns can be added when needed

Column family databases - structure


Column family



- **Parts** of the columns:
 - Name, value, and timestamp

Column family databases - structure

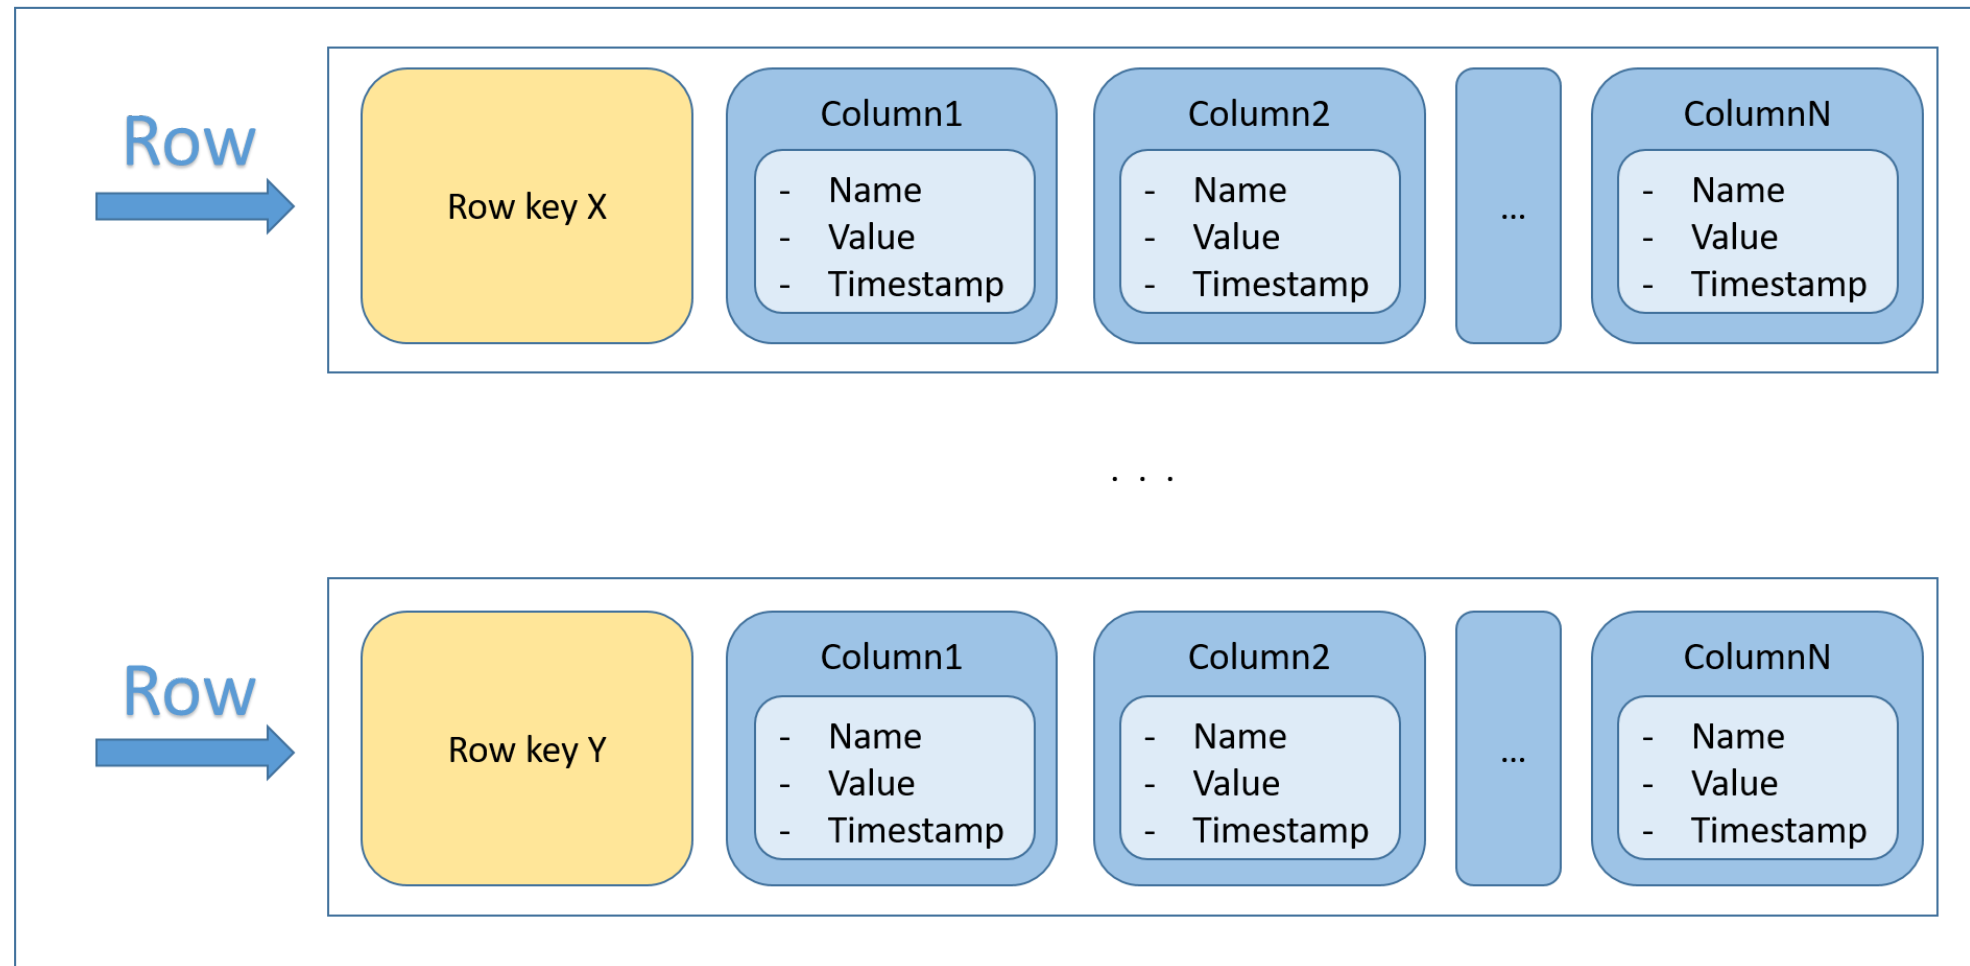
Column family



- **Value:** specify the type depending on the database

Column family databases - structure

Column family



- **Timestamps:** store date and time when the data was inserted.
 - Multiple values of a column

Column family databases - example

512	first name Carol	last name Harper	email carol@datazy.com	address 123 Sesame St, NY
513	first name Benjamin	last name Lieberman	email ben@datazy.com	
514	first name Peter	last name Stallings	email peter@datazy.com	date of birth 07/04/1984

Column family databases - designing

- Think about the queries
- No joins
 - Add all the columns we need

Popular column family databases



Let's practice!
NOSQL CONCEPTS

Advantages and limitations of column family databases

NOSQL CONCEPTS



Miriam Antona
Software engineer

Advantages - flexibility

- Rows within a column family can have **different columns**
- **Add new columns** to a row if we need them
- **Avoids** filling with **default values**
- Flexibility mustn't be considered as the only criterion
 - Evaluate key-value and document databases

Advantages - speed

- Related columns are **stored together on disk**
- **Very fast** writing / retrieving

Advantages - scalability

- Scale horizontally
 - Sharding across multiple servers

Advantages - large volumes of data

- Designed to handle **large volumes of data**
 - speed
 - horizontal scalability
 - efficient data compression

Limitations

- Atomic reads/writes but **no multirow transactions**
- **No joins** support
- **No subqueries** support
- Need to **define the queries** quite well
 - Queries change -> may need to change the column families
 - Can be costly

Let's practice!
NOSQL CONCEPTS

When to use column family databases

NOSQL CONCEPTS




Miriam Antona
Software engineer

Suitable cases - general cases

- Large volumes of data
- Extreme write speeds

Suitable cases - event logging

- Types of events:
 - User logging
 - Errors
 - ...




Suitable cases - Content Management Systems

- Comments
- Links
- Tags
- ...



Suitable cases - time-series data

- Weather
- Traffic
- etc.



Unsuitable cases

- **Prototyping** and at the **beginning** of a project
 - Need to change the queries very frequently
 - Changing the queries -> may imply changing the design of the column families
 - Costly and may slow down the productivity
- Complex queries and joins
- Not dealing with large amounts of data

Let's practice!
NOSQL CONCEPTS


Apache Cassandra case study

NOSQL CONCEPTS



Miriam Antona
Software engineer

Apache Cassandra - overview



Apache

CASSANDRA

- Popular **column family** database
- Originally developed by **Facebook**
- **Open-source**
- Finally became a project of the **Apache Foundation**

Apache Cassandra - features

- **Distributed**
 - Data is distributed across the nodes of the cluster
 - Every node plays the same role
 - No master node
- **High availability**
- **No single point of failure**
- **Scales horizontally** by adding nodes
- Cassandra client **drivers**: C#, Java, Python, Scala, etc.

Apache Cassandra - features

- Cassandra Query Language (aka **CQL**)
 - Query data
 - Similar syntax to SQL
 - Tables (for column families), rows, and columns
 - Differences between CQL and SQL:
 - no joins
 - no foreign keys
 - no subqueries, etc.
 - rows can contain a different number of columns

```
SELECT * FROM users WHERE user_id IN (212, 213, 214);
```

Apache Cassandra - ecosystem

- **Third-party** Cassandra projects, tools, products, and services
 - Cloud offerings
 - Installation tools
 - Developers' frameworks
 - Connectors
 - etc.

Apache Cassandra - customers

NETFLIX



Bigmate case study - overview

- Location tracking
- Industrial sensor
- Productivity



Bigmate case study - problem and solution

- **IoT platform:**
 - Ingests and processes **large volumes** of **different data**
 - **Integrate** IoT sensors, devices, and other platforms
 - Process data in **real-time**
 - Scale and deploy across **multiple locations**
 - Application examples:
 - Thermy -> capture the skin temperature of people
 - Warny -> detects possible collisions
- Tested MySQL, MongoDB, Apache Cassandra, etc.
 - Chose **Apache Cassandra**
 - **Scaled better**

Bigmate case study - results

- **Millions of operations of concurrent users**
- Display 20,000 real-time data points to a single customer
- Fault tolerance (data replication)

Bigmate case study - results

- **Millions of operations of concurrent users**
- Display 20,000 real-time data points to a single customer
- Fault tolerance (data replication)

¹ <https://cassandra.apache.org/case-studies/>

Let's practice!
NOSQL CONCEPTS